



Machine Predictive Maintenance Classification

Names:	ID number:	Task distribution:
Ruyuf alharbi	44510064	<p>Leader/Supervisor.</p> <p>Creating&Editing the report.</p> <p>Writing the introduction, the goal, and the conclusion.</p> <p>Applied three algorithms:</p> <ul style="list-style-type: none"> 1- Randomforset 2- Naïve bayes 3- K- nearest
Alaa Abdulrahman ajlan	445013943	<p>Cleaned the data.</p> <p>Applied three algorithms:</p> <ul style="list-style-type: none"> 1- Logisitic regression. 2- OneR. 3- Decision tree.
Dana Hassan Bafaqih	445013918	<p>Cleaned the data.</p> <p>Applied three algorithms:</p> <ul style="list-style-type: none"> 1- DecisionStump. 2- AdaBoost. 3- DecisionTable
Hadeel Mohammad Turki	445006064	<p>Created the accuracy graph.</p> <p>Applied three algorithms:</p> <ul style="list-style-type: none"> 1- REP tree. 2- Bayes Net. 3- SimpleLogistic.
Refan Salem Almfarriji	445000906	<p>Writing about the ZeroR.</p> <p>Applied three algorithms:</p> <ul style="list-style-type: none"> 1- SMO. 2- Lmt.

Table of content

The introduction:	4
The goal:	4
1) The original dataset information:	5
2) Data cleaning preparation:	7
Machine learning algorithms and methods:	13
1) ZeroR – Zero Rule Classifier	13
2) SMO - Sequential Minimal Optimization	14
3) JRip – Java Repeated Incremental Pruning	16
4) LMT – Logistic Model Tree	17
5) DecissionStump:	19
6) AdaBoostM1:	23
7) DecisionTable:	28
8) Logistic Regression (LR)	35
9) (J48) Tree :	40
10) OneR:	46
11) Simple logistic	50
12) Bayes Network	53
13) REP tree	56
14) Random Forest:	60
15) Naive Bayes	62
14) IBk (K-Nearest Neighbors)	65
The graph	67
<i>A comparison of all the algorithms with the best accuracy</i>	67
The conclusion:	67
References:	68

The introduction:

Predicting when a machine will fail is important to avoid costly repairs and downtime. With data from sensors like temperature, vibration, and pressure, we can use machine learning to spot signs of failure before they actually happen. In This project each one tested three different algorithms to see which one works best at predicting machine failures, helping to keep machines running smoothly and reduce unexpected problems.

The goal:

The goal of this project is to find the best algorithm to predict if a machine will fail or not. We cleaned the data and wanted to build a model that can clearly tell which machines might fail. This helps plan maintenance early and avoid surprise breakdowns.

1) The original dataset information:

Machine Predictive Maintenance Classification Dataset

Real predictive maintenance datasets are hard to get and share, so this dataset is synthetic but designed to closely reflect real industrial scenarios.

- The dataset has **10,000 rows** (data points) and **14 features** (columns).

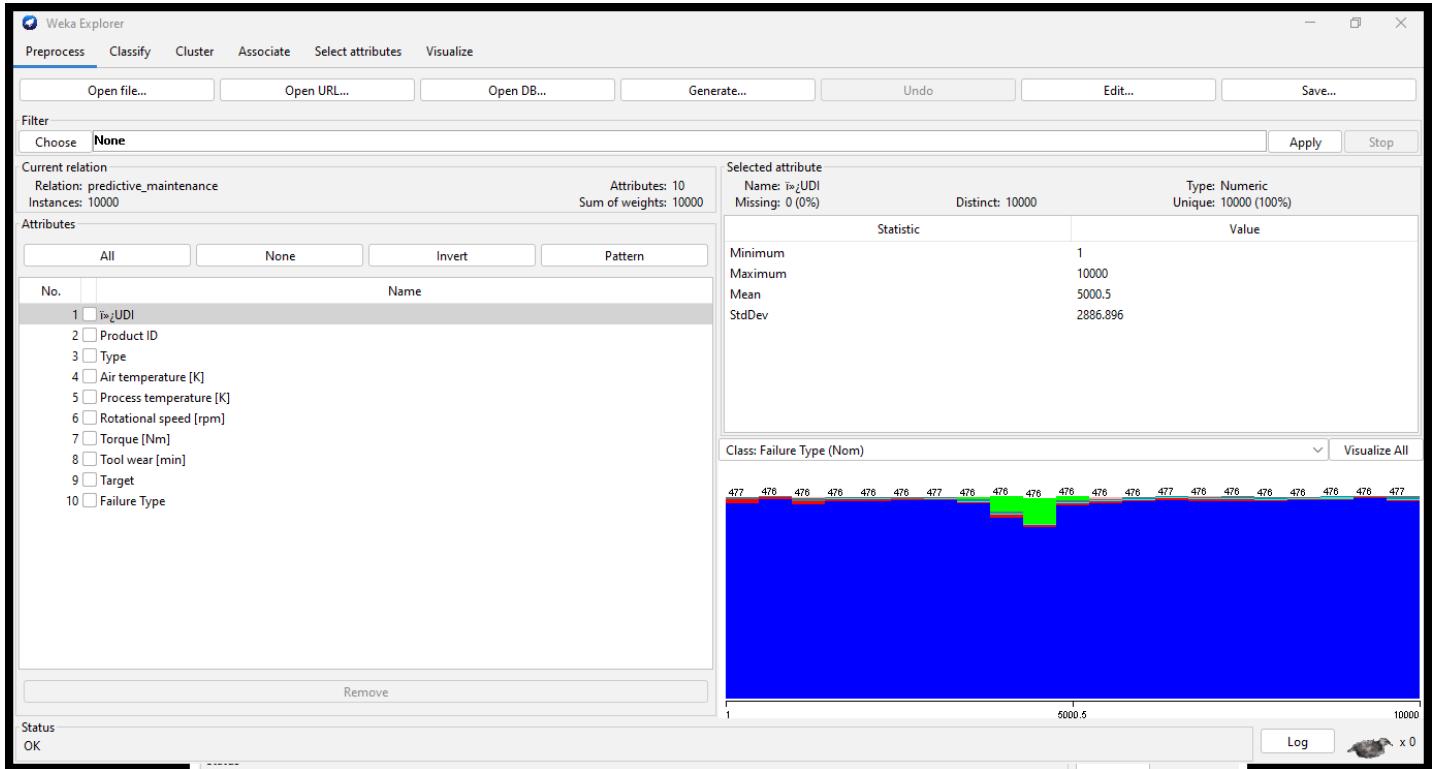
Features include:

- **UID**: Unique identifier from 1 to 10,000.
- **productID**: Product quality variant labeled as L (low, 50%), M (medium, 30%), or H (high, 20%), followed by a serial number.
- **Air temperature [K]**: Generated by a random walk and normalized with a standard deviation of 2 K around 300 K.
- **Process temperature [K]**: Based on air temperature plus 10 K, with additional noise normalized at 1 K.
- **Rotational speed [rpm]**: Calculated from power of 2860 W with added normal noise.
- **Torque [Nm]**: Normally distributed around 40 Nm ($\sigma = 10$ Nm), no negative values.
- **Tool wear [min]**: Quality variants add 5, 3, or 2 minutes of tool wear for H, M, and L respectively.
- **Machine failure label**: Indicates if the machine failed at this data point due to any failure mode.

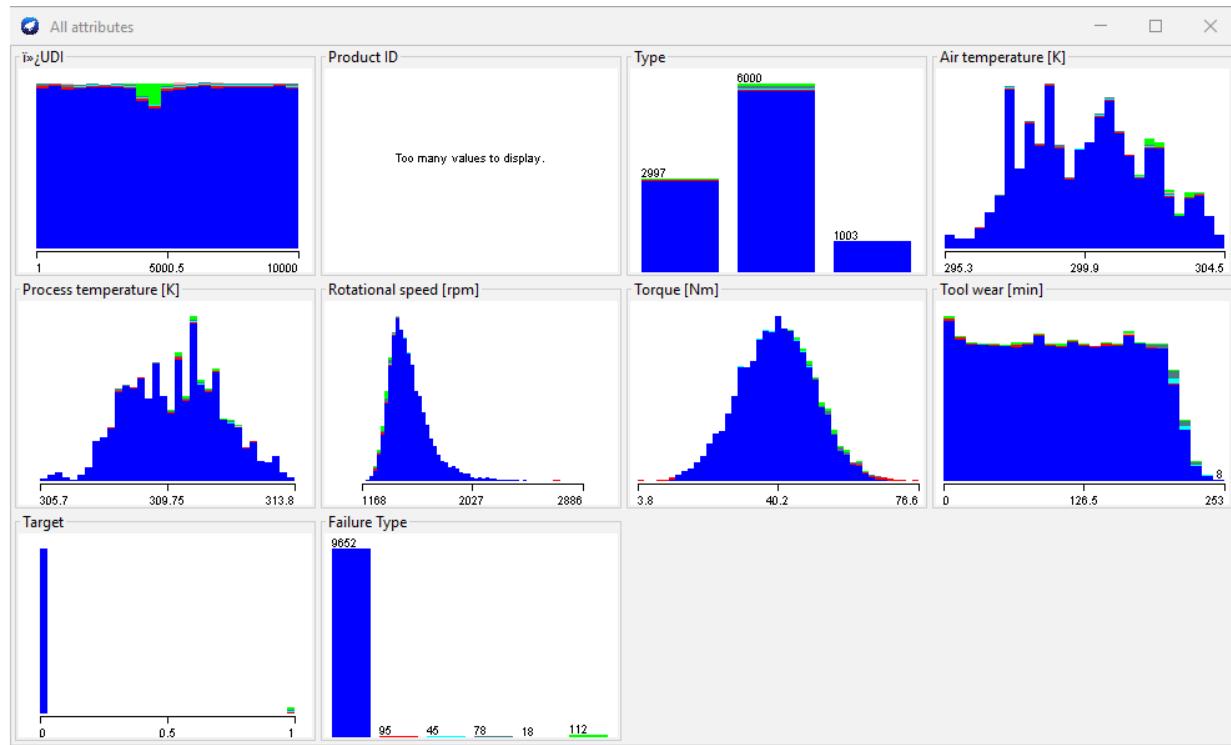
Important: The dataset contains **two target variables**:

- **Target**: Whether the machine failed or not.
- **Failure Type**: The type of failure.

- The dataset



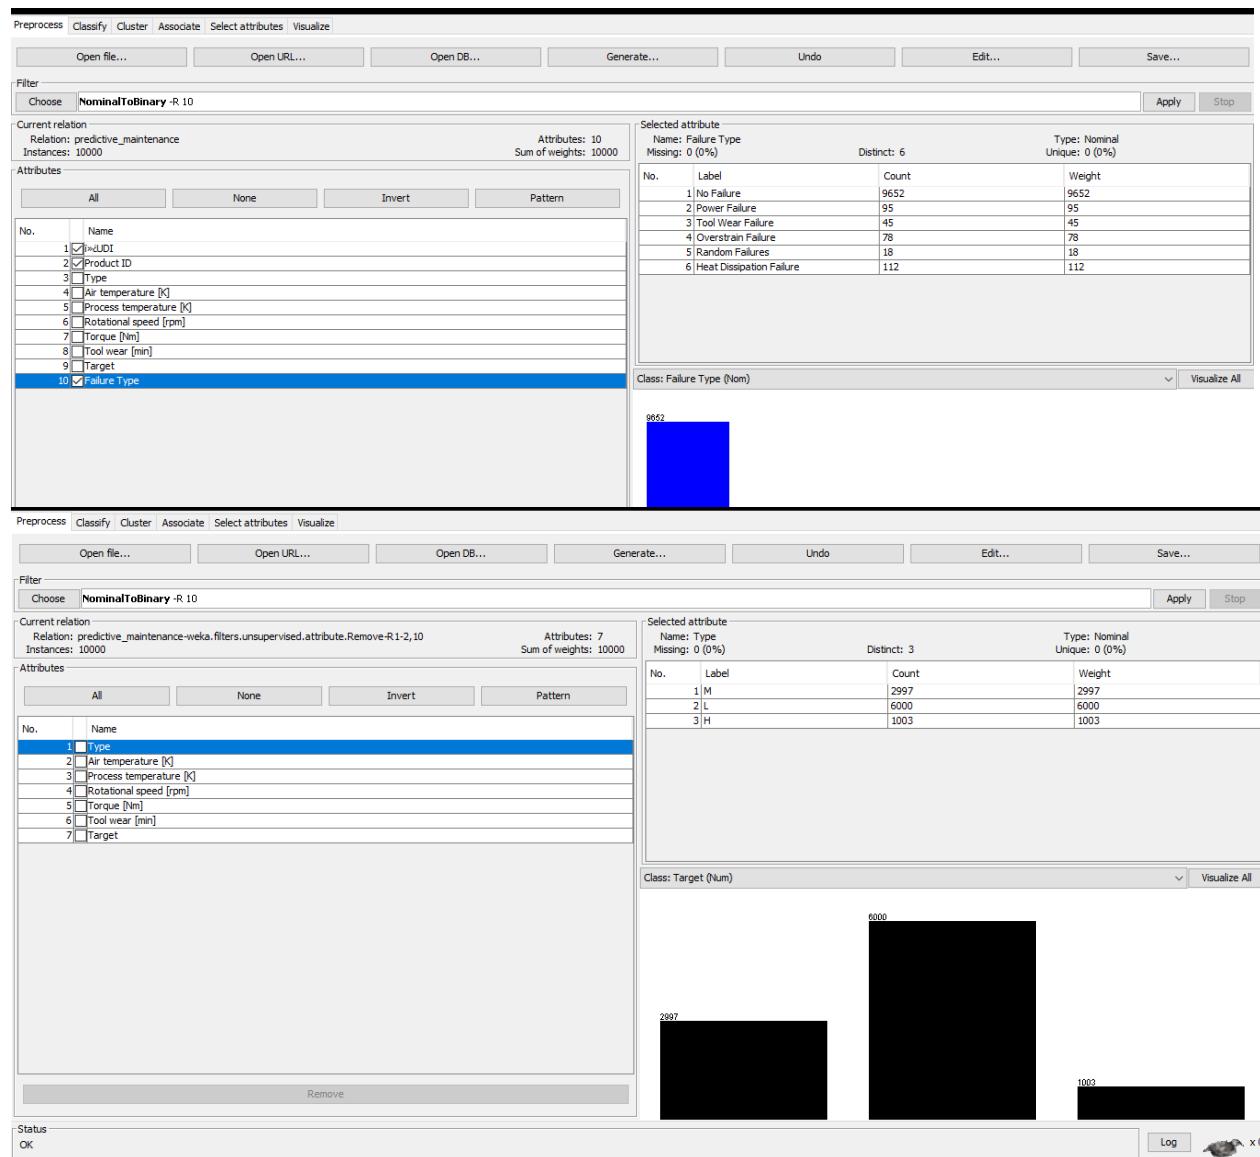
- Visualization



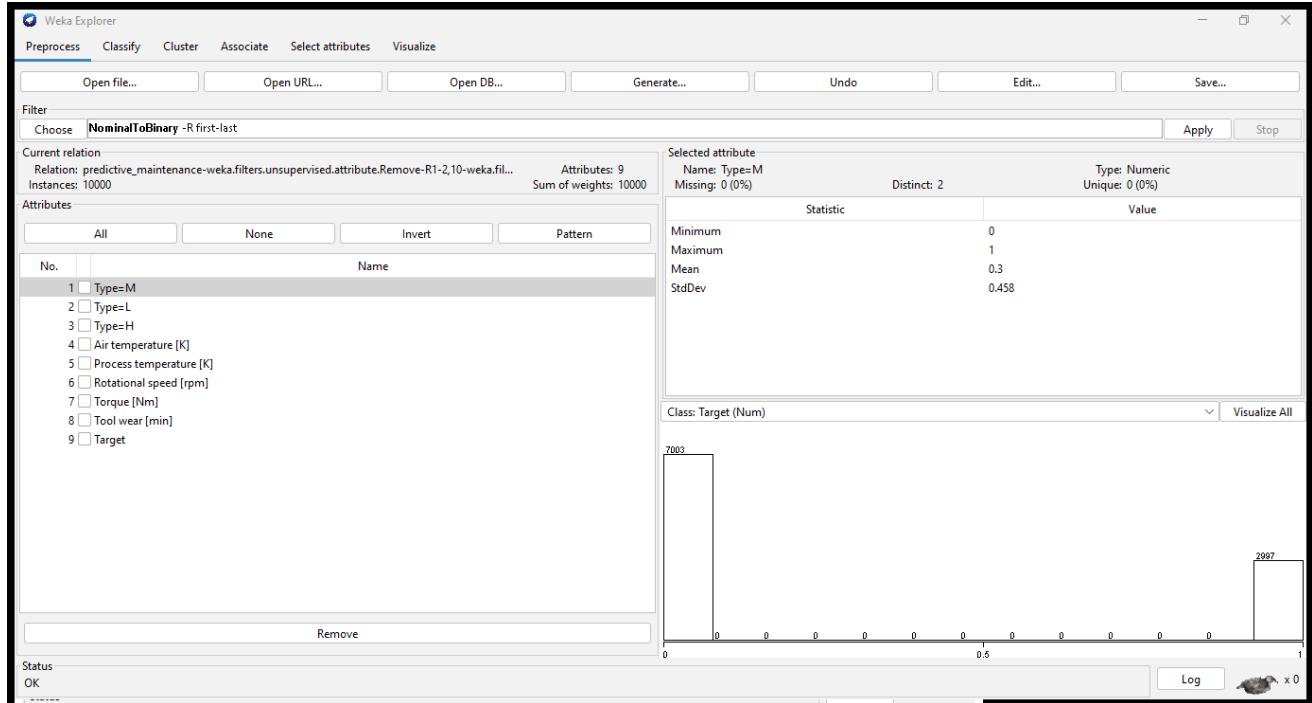
2) Data cleaning preparation:

We applied some filter in order to clean and prepare the dataset for model training and evaluation:

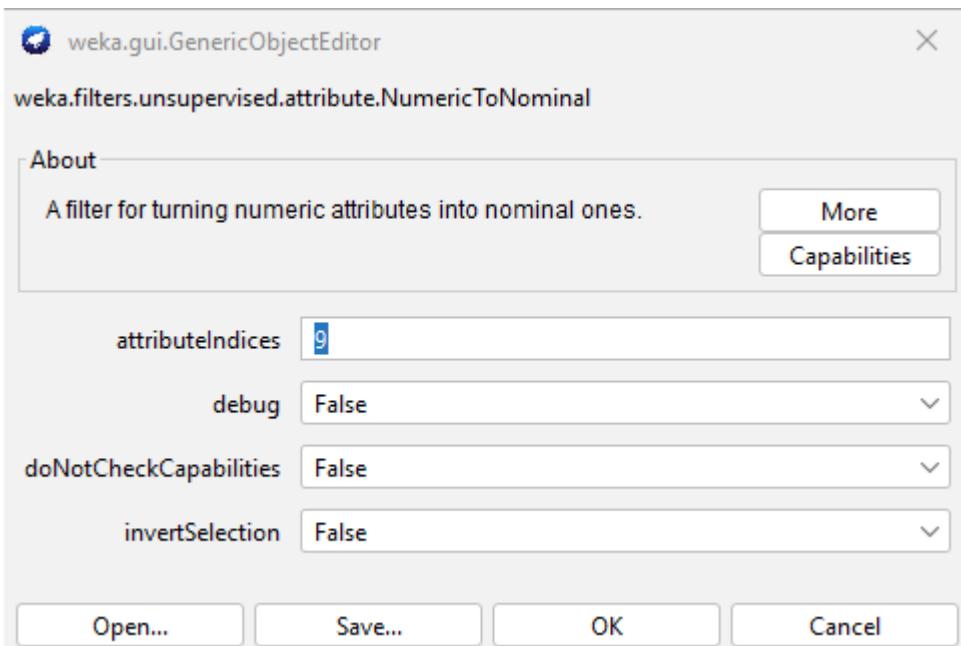
- First, we removed the columns (UID, Product ID) that were not useful for the prediction process. This helped improve the model's performance and reduce complexity.
- We also removed the (failure type)column because we want to predict whether the machine will fail or not, and that column does not help with the prediction.

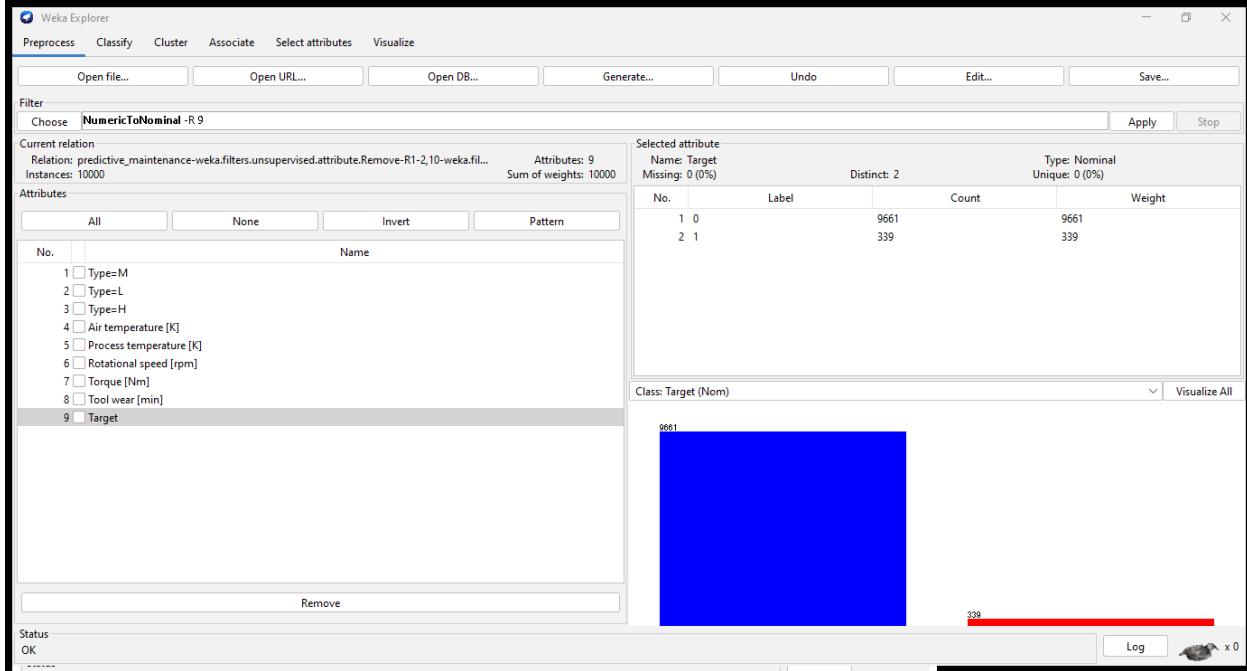


- We converted the "Type" column to binary to make it easier to handle. Additionally, some algorithms cannot work with categorical values directly and require numerical input for processing.

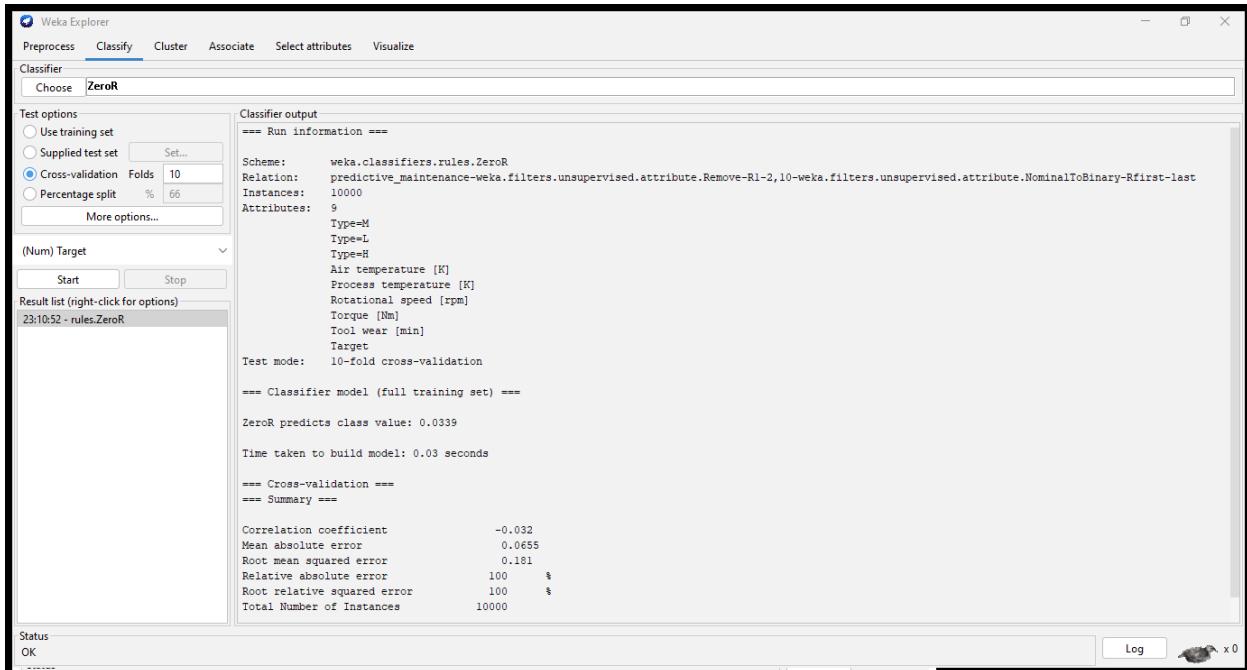


- The target column had numerical values, and since we want to use it for classification, we converted it to nominal.
 - We chose only the target column

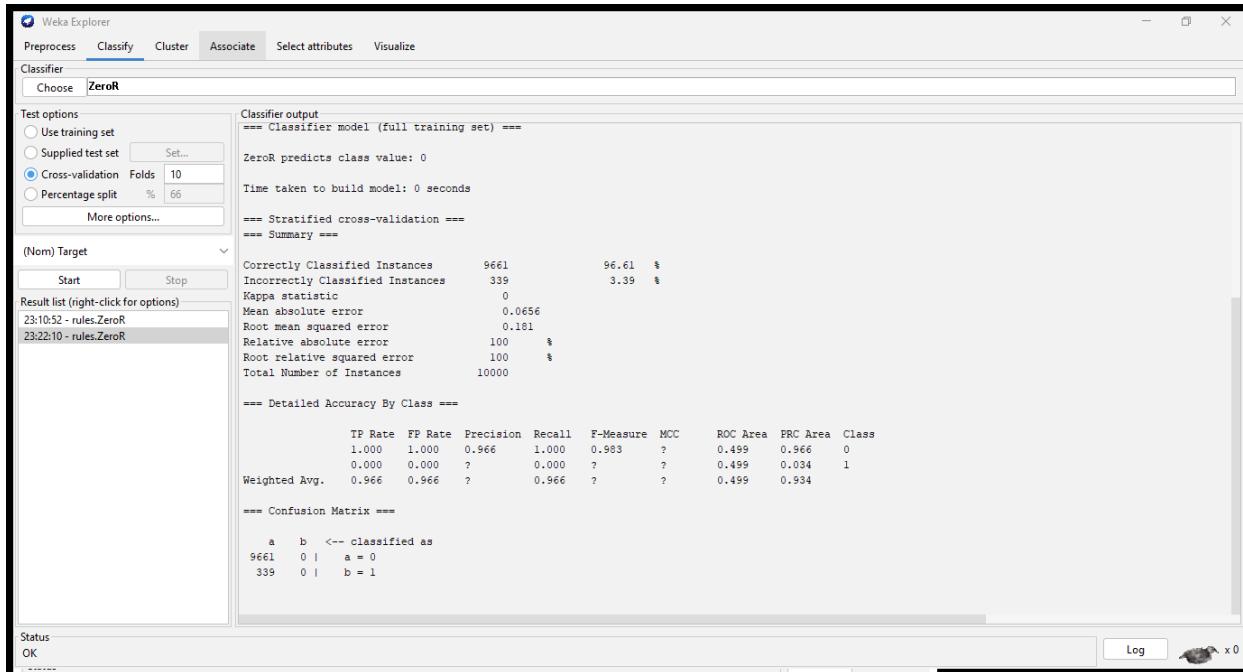




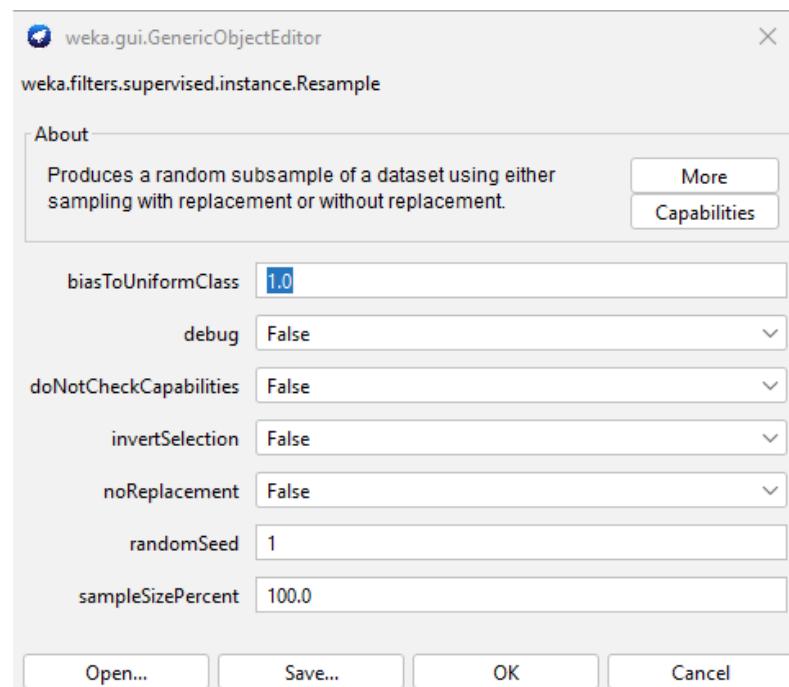
- The result before converting the column was treated as a regression problem, so the model predicted a number instead of a class.

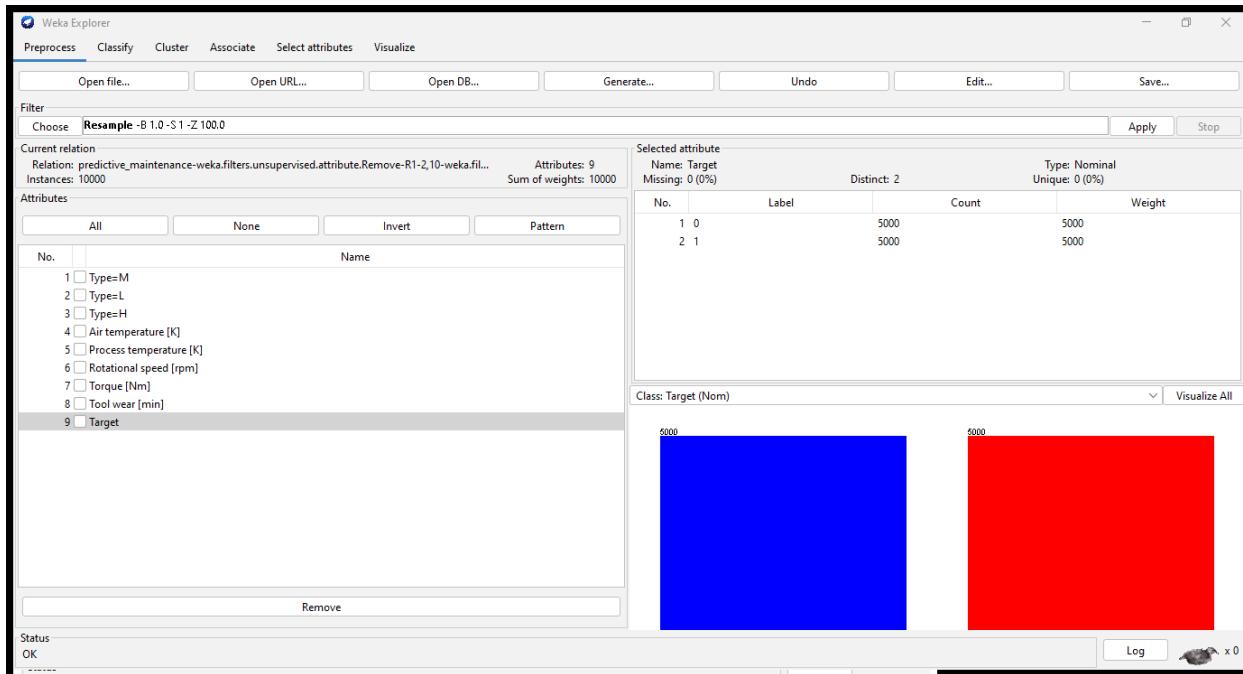


- The result after the conversion

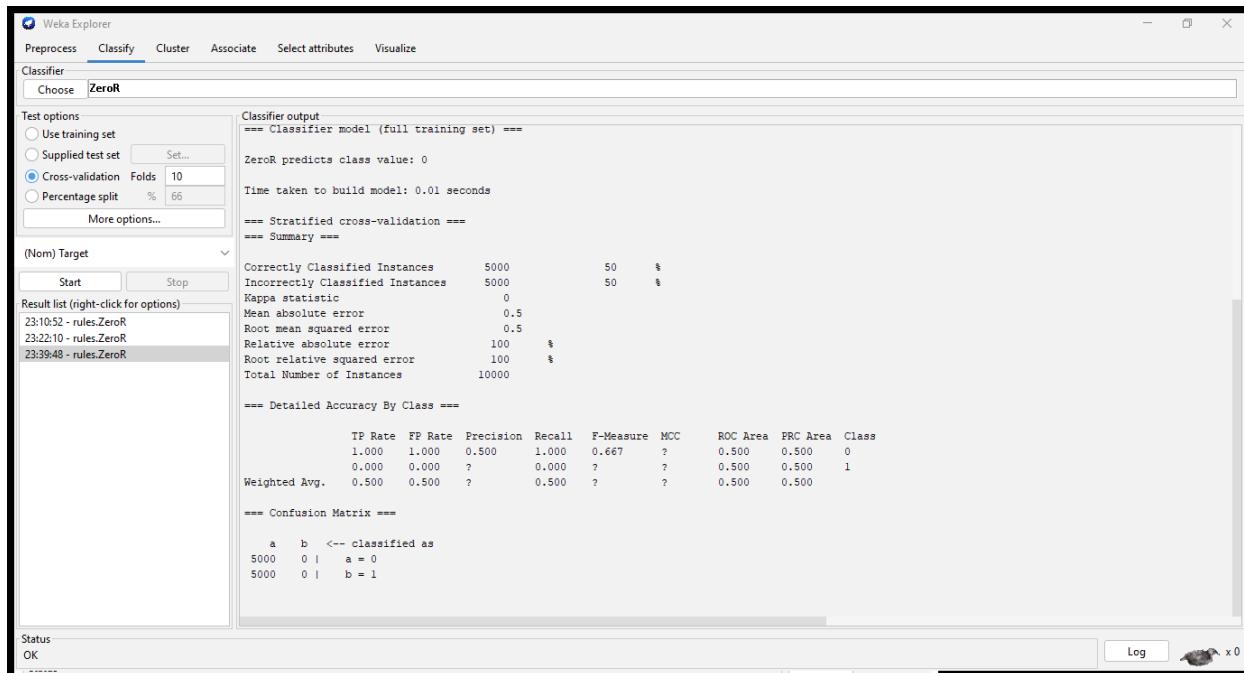


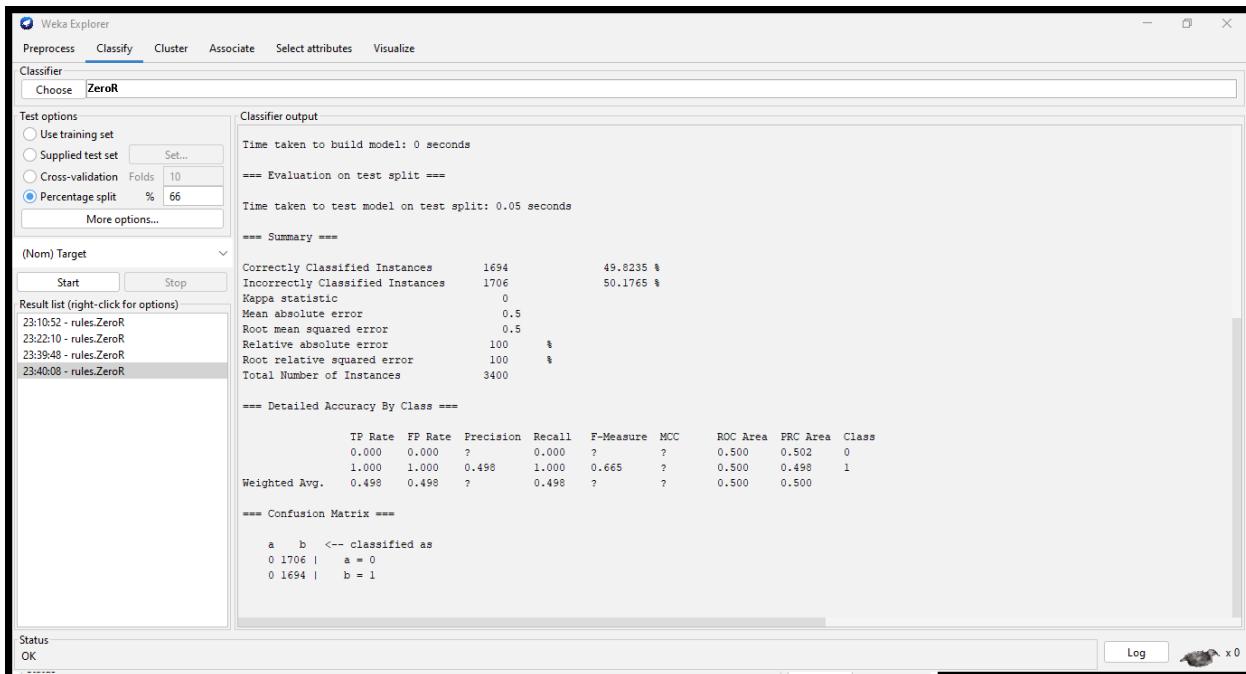
- As we can see, our data is imbalanced, and this causes a problem in classification. To solve this, we used the **Resample filter**, which helps make the data more balanced.
 - **Bias to Uniform Class = 1**, which means it will try to make the class distribution equal.
 - **Sample Size = 100%**, so the number of instances stays the same.
 - **No Replacement = false**, which means Weka is allowed to repeat some data.





- The result after the cleaning:





We did not perform normalization or discretization on the dataset manually, because during our research we found that these preprocessing steps can be applied directly through the algorithm's parameter settings in Weka if needed. Therefore, we relied on the algorithm's internal options instead of applying them externally.

Machine learning algorithms and methods:

1) ZeroR – Zero Rule Classifier

ZeroR is the simplest classification algorithm. It ignores all input features and always predicts the majority class in the training data. It is mainly used as a baseline model to evaluate the performance of more advanced classifiers.

Experimental Result:

In our experiment, ZeroR predicted only class 0 for all instances. As a result, it achieved exactly 50% accuracy, since the dataset was balanced (50% class 0, 50% class 1). The confusion matrix confirmed that all 5000 instances of class 0 were classified correctly, while all 5000 instances of class 1 were misclassified.

The screenshot shows the Weka GUI interface. At the top, a window titled "weka.gui.GenericObjectEditor" displays the "weka.classifiers.rules.ZeroR" class information, stating it's a "Class for building and using a 0-R classifier". Below this, several configuration parameters are set: batchSize to 100, debug to False, doNotCheckCapabilities to False, and numDecimalPlaces to 2. At the bottom of this window are "Open...", "Save...", "OK", and "Cancel" buttons. The main Weka window below has tabs for Preprocess, Classify, Cluster, Associate, Selectattributes, and Visualize. The "Classifier" tab is selected, showing "Choose ZeroR". Under "Test options", "Cross-validation" is selected with "Folds" set to 10. The "Classifier output" pane displays the following text:

```
Torque [Nm]
Tool wear [min]
Target
Test mode: 10-fold cross-validation
==== Classifier model (full training set) ====
ZeroR predicts class value: 0
Time taken to build model: 0 seconds
==== Stratified cross-validation ====
==== Summary ===
Correctly Classified Instances      5000      50      %
Incorrectly Classified Instances   5000      50      %
Kappa statistic                   0
Mean absolute error               0.5
Root mean squared error           0.5
Relative absolute error            100      %
Root relative squared error       100      %
Total Number of Instances         10000
==== Detailed Accuracy By Class ====

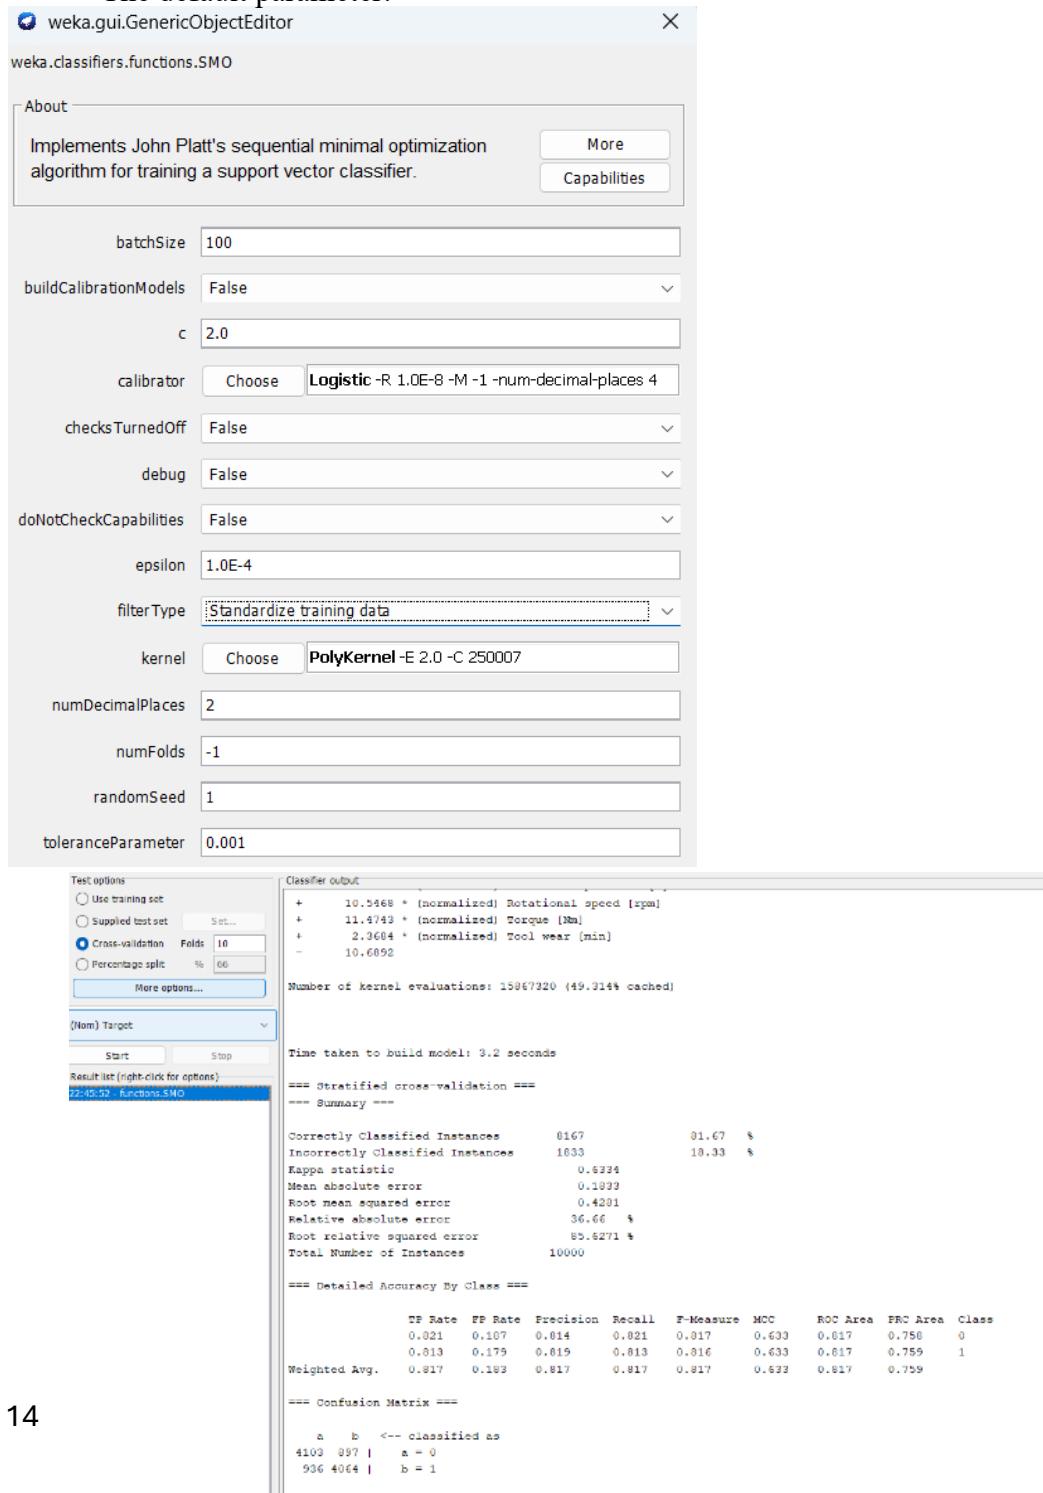
TP Rate  FP Rate  Precision  Recall  F-Measure  MCC    ROC Area  PRC Area  Class
1.000   1.000    0.500    1.000   0.667    ?     0.500   0.500    0
0.000   0.000    ?        0.000   ?        ?     0.500   0.500    1
Weighted Avg.                    0.500   0.500    ?        0.500   ?        ?     0.500   0.500
==== Confusion Matrix ====

a    b    <-- classified as
5000  0 |  a = 0
5000  0 |  b = 1
```

2) SMO - Sequential Minimal Optimization

SMO trains an SVM model by breaking down the large optimization problem into a series of smaller problems involving only two variables at a time. Each small problem is solved analytically, without the need for complex iterative methods. This process is repeated until the entire model is trained efficiently

- The default parameter:



Best Accuracy Result:

SMO achieved its highest accuracy (**93.95%**) after changing the kernel to *PUK*, applying the **Normalize filter** to the training data, and **increasing the number of folds**

The screenshot shows the Weka GUI interface for the SMO classifier. On the left, the configuration panel displays various parameters such as batchSize (100), buildCalibrationModels (False), c (2.0), calibrator (Logistic -R 1.0E-8 -M 1 -num-decimal-places 4), checksTurnedOff (False), debug (False), doNotCheckCapabilities (False), epsilon (1.0E-4), filterType (Normalize training data), kernel (Puk -O 1.0 -S 1.0 -C 250007), numDecimalPlaces (2), numFolds (5), randomSeed (1), and toleranceParameter (0.001). On the right, the classifier output window shows the results of a 10-fold cross-validation. The output includes the command used, test options, classifier settings, and detailed performance metrics. The accuracy is summarized as 93.95% correctly classified instances and 6.05% incorrectly classified instances. A confusion matrix is also provided.

```

Classifier
Choose SMO -C 0.2 -L 0.001 -P 1.0E-4 -N 0 -V 5 -W 1 -K "weka.classifiers.functions.supportVector.Puk -O 1.0 -S 1.0 -C 250007" -calibrator "weka.classifiers.functions.Logistic -R 1.0E-8 -M 1 -num-decimal-places 4"

Test options
 Use training set
 Supplied test set Set...
 Cross-validation Folds 10
 Percentage split % 66
More options...

Classifier output
+ 0.0966 * <= 0 1 0.068132 0.7 0.105355 0.649725 0.035573 > * x]
+ 1.8375

Number of support vectors: 3926
Number of kernel evaluations: 82913086 (16.208% cached)

(Nom) Target
Start Stop
Result list (right-click for options)
22:49:52 - functions.SMO
22:57:06 - functions.SMO
23:06:39 - functions.SMO
23:11:41 - functions.SMO
23:16:12 - functions.SMO
23:16:12 - functions.SMO

Time taken to build model: 10.74 seconds
*** Stratified cross-validation ===
*** Summary ===

Correctly Classified Instances 9395 93.95 %
Incorrectly Classified Instances 605 6.05 %
Kappa statistic 0.879
Mean absolute error 0.0605
Root mean squared error 0.246
Relative absolute error 12.1 %
Root relative squared error 49.1935 %
Total Number of Instances 10000

*** Detailed Accuracy By Class ===

TP Rate FP Rate Precision Recall F-Measure MCC ROC Area PRC Area Class
0.914 0.035 0.963 0.914 0.938 0.880 0.940 0.923 0
0.965 0.006 0.918 0.965 0.941 0.880 0.940 0.904 1

Weighted Avg. 0.940 0.061 0.941 0.940 0.939 0.880 0.940 0.913 0.913

*** Confusion Matrix ===

a b <-> classified as
4570 430 | a = 0
175 4025 | b = 1

```



```

Classifier
Choose SMO -C 0.2 -L 0.001 -P 1.0E-4 -N 0 -V 5 -W 1 -K "weka.classifiers.functions.supportVector.Puk -O 1.0 -S 1.0 -C 250007" -calibrator "weka.classifiers.functions.Logistic -R 1.0E-8 -M 1 -num-decimal-places 4"

Test options
 Use training set
 Supplied test set Set...
 Cross-validation Folds 10
 Percentage split % 66
More options...

Classifier output
Number of support vectors: 3926
Number of kernel evaluations: 82913086 (16.208% cached)

Time taken to build model: 12.37 seconds
*** Evaluation on test split ===

Time taken to test model on test split: 0.74 seconds
*** Summary ===

Correctly Classified Instances 3189 93.7941 %
Incorrectly Classified Instances 211 6.2059 %
Kappa statistic 0.8759
Mean absolute error 0.0621
Root mean squared error 0.2491
Relative absolute error 12.4117 %
Root relative squared error 49.8228 %
Total Number of Instances 3400

*** Detailed Accuracy By Class ===

TP Rate FP Rate Precision Recall F-Measure MCC ROC Area PRC Area Class
0.899 0.022 0.976 0.899 0.936 0.879 0.938 0.928 0
0.978 0.101 0.905 0.978 0.940 0.879 0.938 0.896 1

Weighted Avg. 0.938 0.062 0.941 0.938 0.938 0.879 0.938 0.912 0.912

*** Confusion Matrix ===

a b <-> classified as
1533 173 | a = 0
39 1656 | b = 1

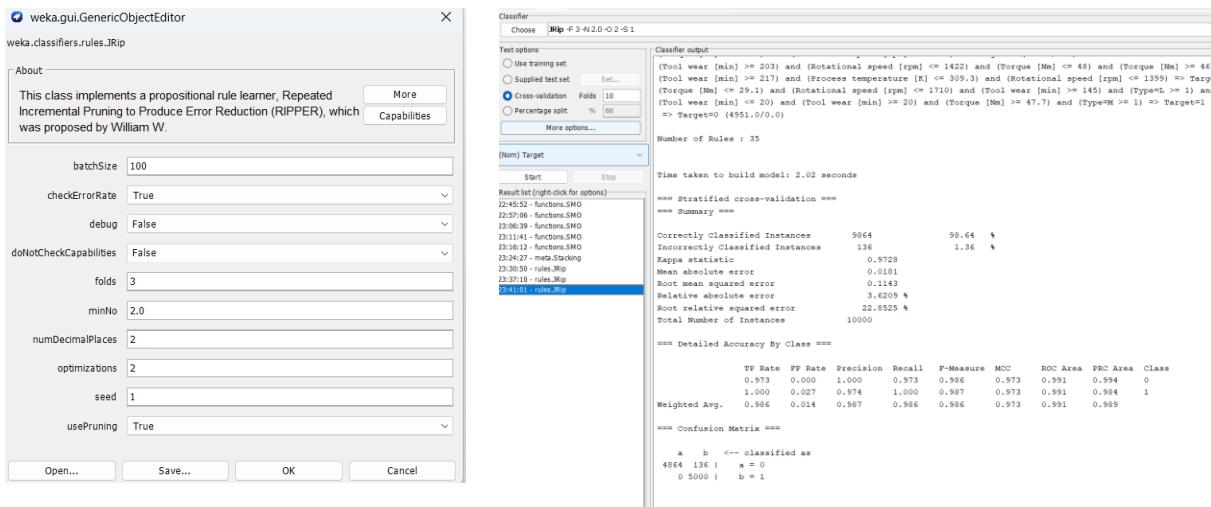
```

3) JRip – Java Repeated Incremental Pruning

JRip is a rule-based classification algorithm that builds simple if-then rules to classify data. It improves accuracy by pruning unnecessary parts of the rules and is useful when clear and interpretable models are needed

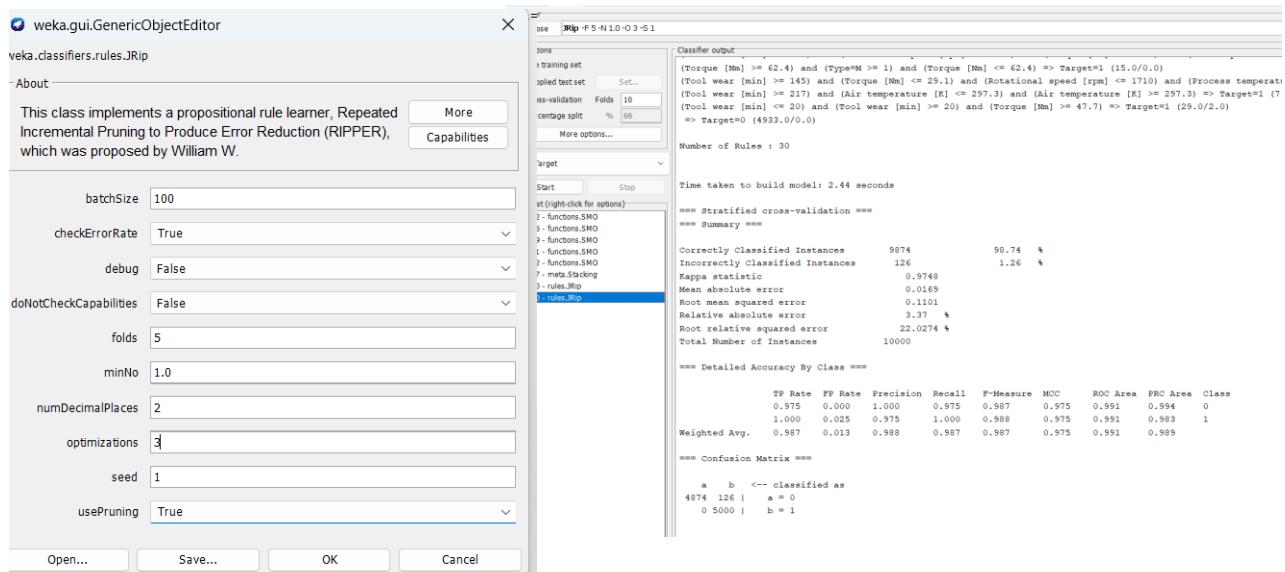
- The default parameter:

The default JRip parameters already achieved high accuracy (98.64%). However, I made a slight adjustment to one of the parameters to further improve the performance



Best Accuracy Result:

JRip achieved its highest accuracy (98.74%) after increasing the optimizations parameter , decreasing number of fold ,and decreasing numNO



Test options

- Use training set
- Supplied testset Set...
- Cross-validation Folds 10
- Percentage split % 66
- [More options...](#)

Classifier output

```
(Tool wear [min] <= 20) and (Tool wear [min] >= 20) and (Torque [Nm] >= 47.7) => Target=0 (4933.0/0.0)
=> Target=0 (4933.0/0.0)
```

Number of Rules : 30

Time taken to build model: 3.61 seconds

==== Evaluation on test split ===

Time taken to test model on test split: 0 seconds

==== Summary ===

	Correctly Classified Instances	3353	98.6176 %
Incorrectly Classified Instances	47	1.3824 %	
Kappa statistic	0.9724		
Mean absolute error	0.0177		
Root mean squared error	0.1165		
Relative absolute error	3.5388 %		
Root relative squared error	23.2965 %		
Total Number of Instances	3400		

==== Detailed Accuracy By Class ===

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	FRC Area	Class
0	0.974	0.002	0.998	0.974	0.986	0.973	0.987	0.987	0
1	0.998	0.026	0.975	0.998	0.986	0.973	0.987	0.975	1
Weighted Avg.	0.986	0.014	0.986	0.986	0.986	0.973	0.987	0.981	

==== Confusion Matrix ===

	a	b	<- classified as
a	1662	44	a = 0
b	3	1691	b = 1

4) LMT – Logistic Model Tree

LMT is a hybrid classification algorithm that combines decision trees and logistic regression. It builds a tree where each leaf node contains a logistic regression model, allowing the algorithm to handle both non-linear and linear patterns in data effectively.

weka.gui.GenericObjectEditor

weka.classifiers.trees.LMT

About

Classifier for building "logistic model trees", which are classification trees with logistic regression functions at the leaves.

Test options

- Use training set
- Supplied testset Set...
- Cross-validation Folds 10
- Percentage split % 66
- [More options...](#)

Classifier Choose **LMT-1-1-M15-W0.0**

Result list (right-click for options)

```
[Type=L] * 3.45 +
[Type=H] * 0.28 +
[Air temperature [K]] * 0 +
[Process temperature [K]] * -0.07 +
[Rotational speed [rpm]] * 0.03 +
[Torque [Nm]] * 0.52 +
[Tool wear [min]] * 0.12
```

Time taken to build model: 3.9 seconds

==== Stratified cross-validation ===

==== Summary ===

	Correctly Classified Instances	9906	99.06 %
Incorrectly Classified Instances	94	0.94 %	
Kappa statistic	0.9812		
Mean absolute error	0.0106		
Root mean squared error	0.0958		
Relative absolute error	2.1153 %		
Root relative squared error	19.1600 %		
Total Number of Instances	10000		

==== Detailed Accuracy By Class ===

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	FRC Area	Class
0	0.981	0.000	1.000	0.981	0.991	0.981	0.994	0.996	0
1	1.000	0.019	0.982	1.000	0.991	0.981	0.994	0.988	1
Weighted Avg.	0.991	0.009	0.991	0.991	0.991	0.981	0.994	0.992	

==== Confusion Matrix ===

	a	b	<- classified as
a	4906	94	a = 0
b	0	5000	b = 1

batchSize 100

convertNominal False

debug False

doNotCheckCapabilities False

doNotMakeSplitPointActualValue False

errorOnProbabilities False

fastRegression True

minNumInstances 15

numBoostingIterations -1

numDecimalPlaces 2

splitOnResiduals False

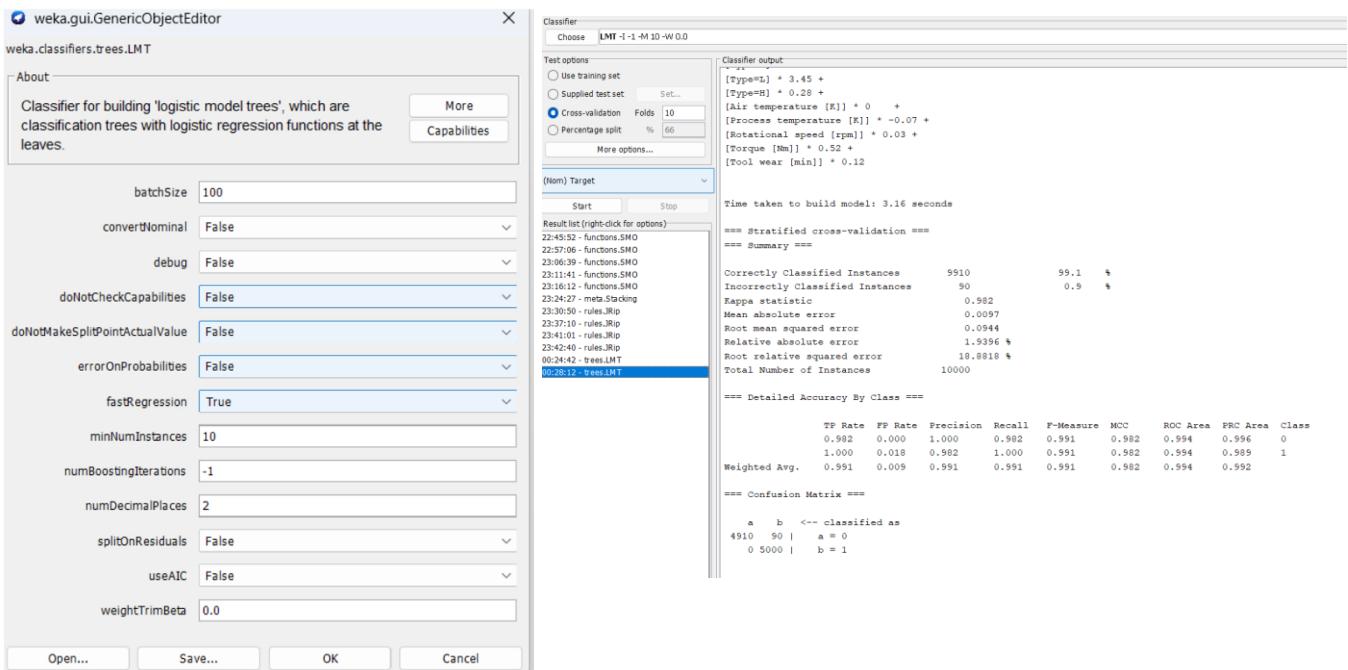
useAIC False

weightTrimBeta 0.0

Open... **Save...** **OK** **Cancel**

Best Accuracy:

This was the best accuracy achieved in all experiments, reaching 99.1% after adjusting the minNumInstances parameter



Classifier output

```

[Process temperature [K]] * -0.07 +
[Rotational speed [rpm]] * 0.03 +
[Torque [Nm]] * 0.52 +
[Tool wear [min]] * 0.12

Time taken to build model: 3.1 seconds
==== Evaluation on test split ====
Time taken to test model on test split: 0 seconds
==== Summary ====
Correctly Classified Instances      3371          99.1471 %
Incorrectly Classified Instances    29           0.8529 %
Kappa statistic                   0.9829
Mean absolute error               0.0098
Root mean squared error           0.0916
Relative absolute error           1.9651 %
Root relative squared error      18.3277 %
Total Number of Instances         3400

==== Detailed Accuracy By Class ====

      TP Rate  FP Rate  Precision  Recall  F-Measure  MCC   ROC Area  PRC Area  Class
          0.983   0.000    1.000    0.983   0.991   0.983   0.994   0.996   0
          1.000   0.017    0.983    1.000    0.992   0.983   0.994   0.990   1
Weighted Avg.        0.991   0.008    0.992    0.991   0.991   0.983   0.994   0.993

==== Confusion Matrix ====

      a     b  <- classified as
  1677   29 |  a = 0
  0 1694 |  b = 1
  
```

5) DecisionStump:

The Decision Stump focuses on only **one feature** at a time and finds a threshold that best separates the data.

In our data, it makes the decision based on the '**Rotational speed**' feature.

```
Decision Stump

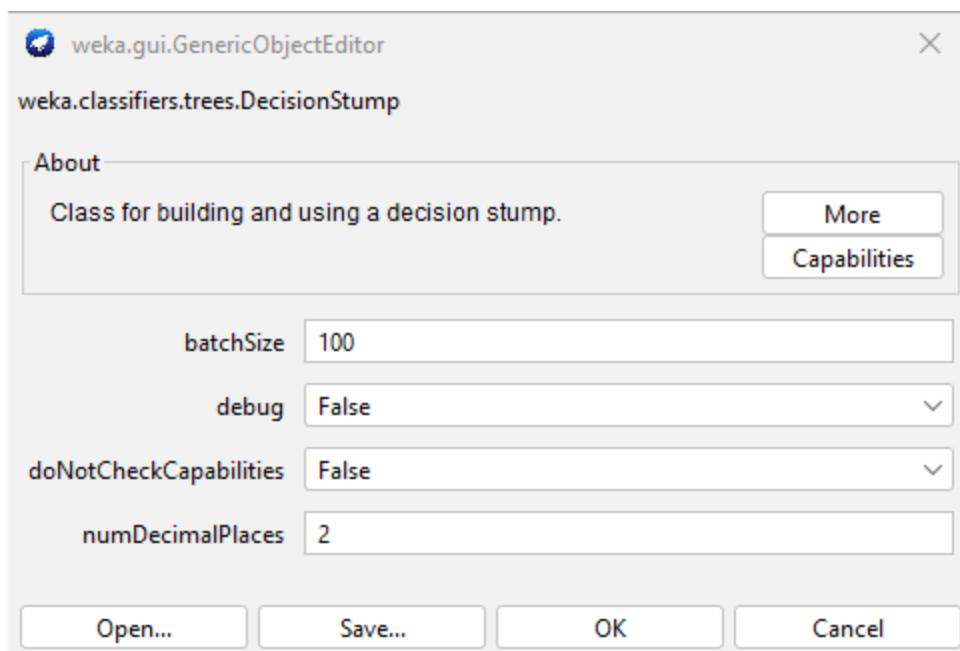
Classifications

Rotational speed [rpm] <= 1381.5 : 1
Rotational speed [rpm] > 1381.5 : 0
Rotational speed [rpm] is missing : 0

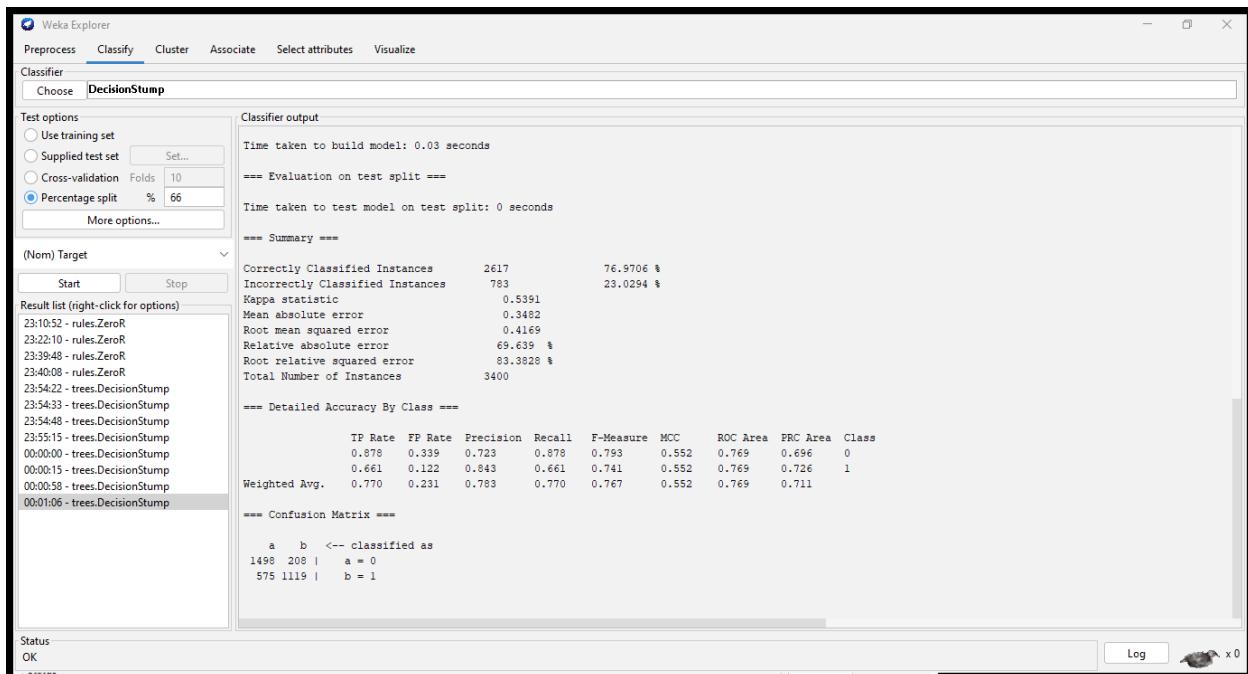
Class distributions

Rotational speed [rpm] <= 1381.5
0      1
0.154446106240330068    0.8455389375966993
Rotational speed [rpm] > 1381.5
0      1
0.7188827180659915    0.2811172819340085
Rotational speed [rpm] is missing
0      1
0.5    0.5
```

- The default parameters:



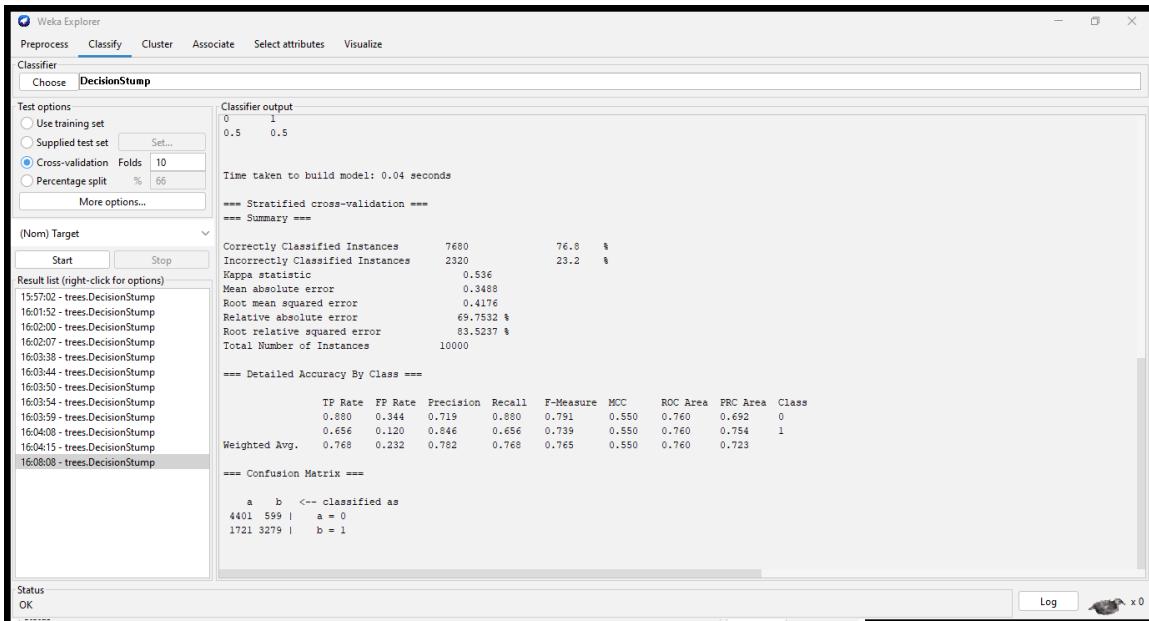
- The result:



The accuracy: 76.97%

The confusion matrix:

	Predicted 0	Predicted 1
Actual 0	1498	208
Actual 1	575	1119



The accuracy: 76.8%

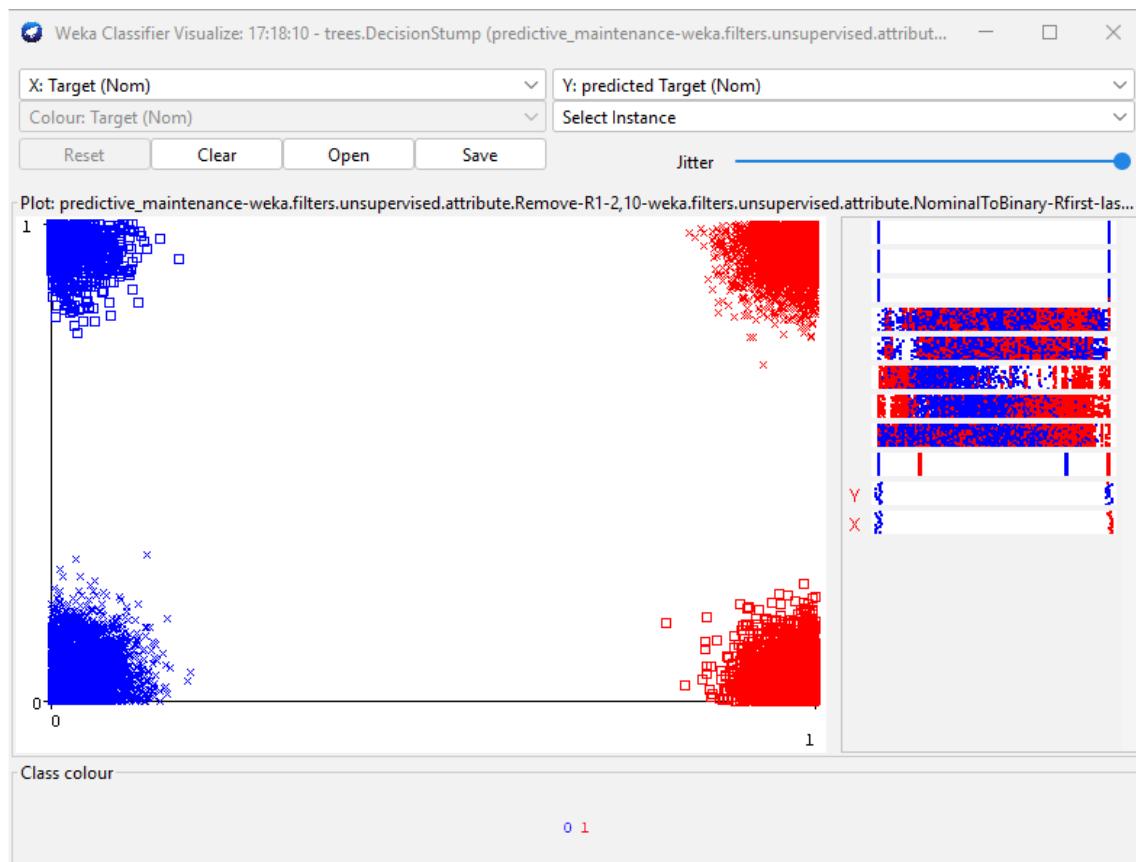
The confusion matrix:

	Predicted 0	Predicted 1
Actual 0	4401	599
Actual 1	1721	3279

- **The algorithm with different parameters:**

- I used the algorithm with different parameters, but the results did **not** change, and the accuracy remained the same. However, the overall accuracy is **not** satisfactory.

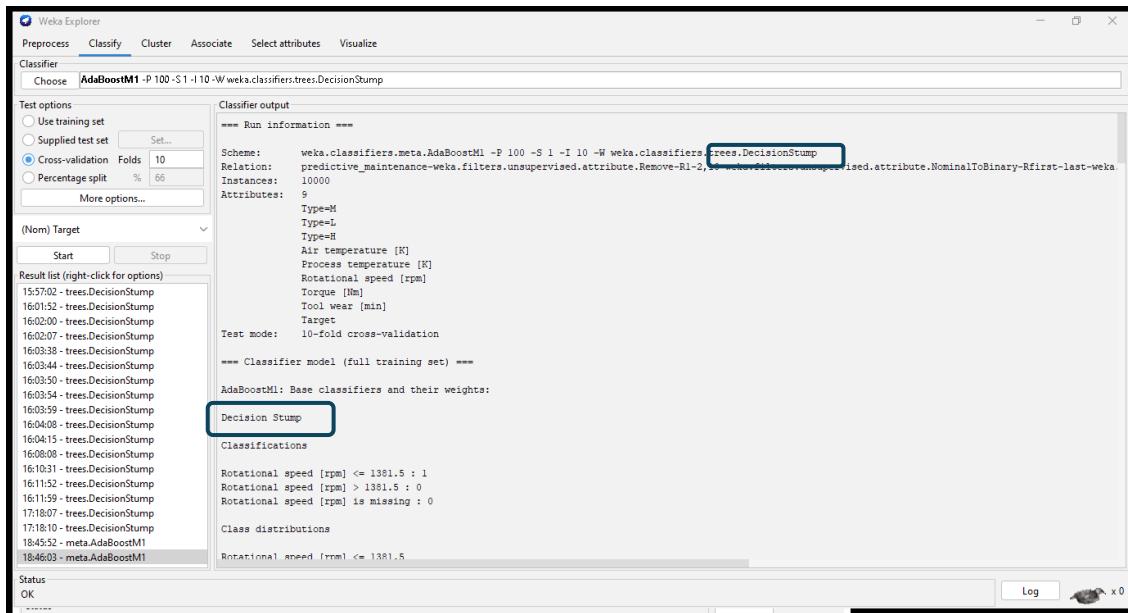
- Visualize classifier Error:



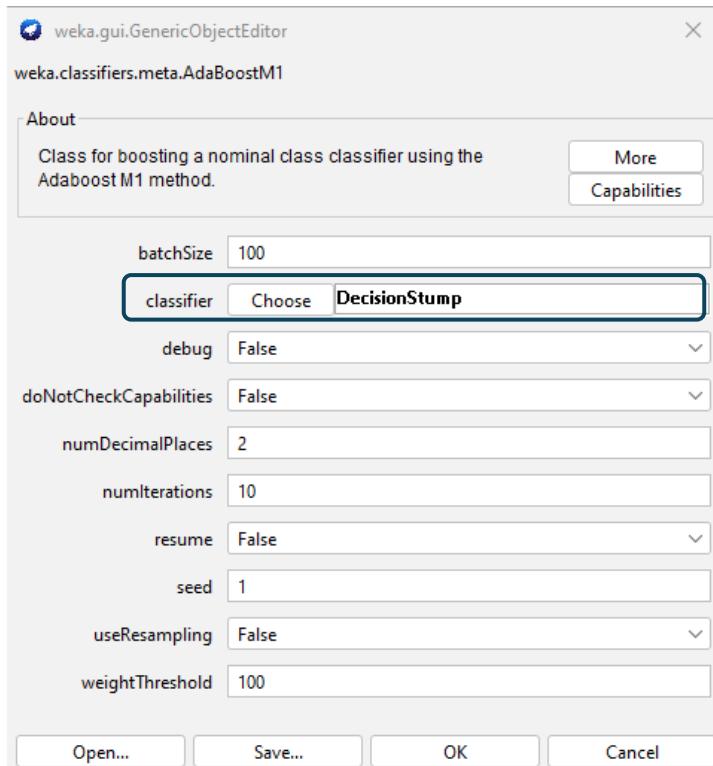
The **x**-axis represent the **true** values, while the **y**-axis represent the **predicted** values

6) AdaBoostM1:

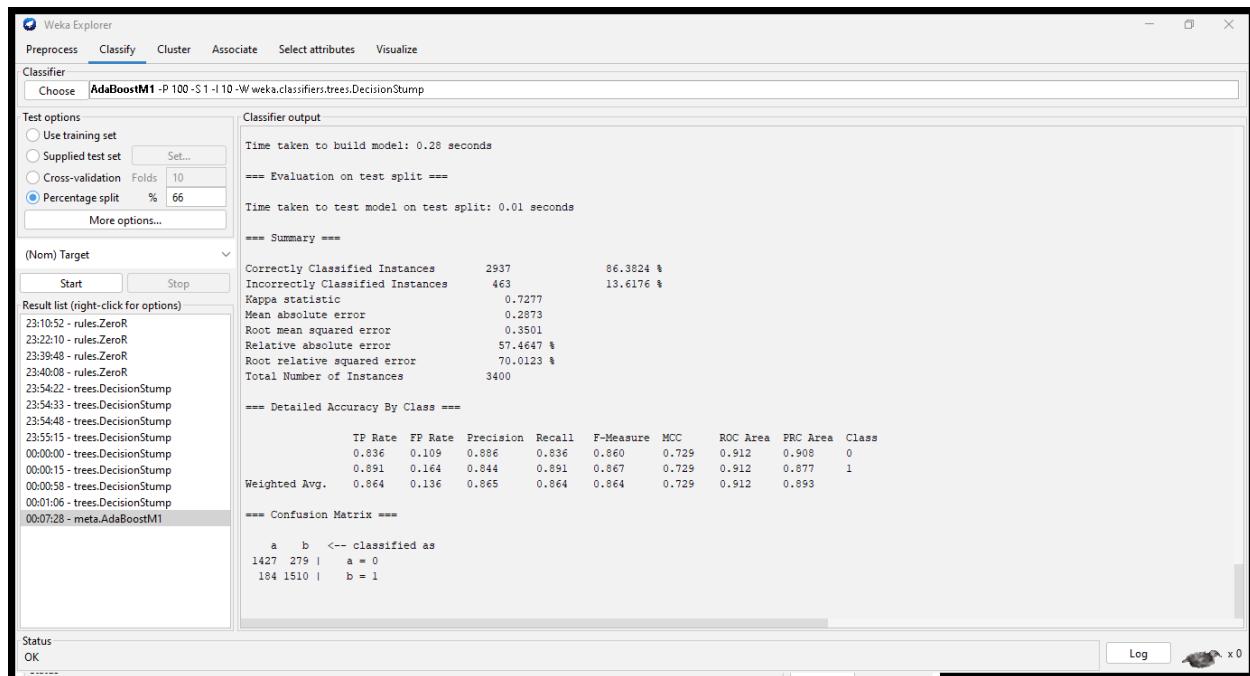
- The AdaBoostM1 combines multiple weak classifiers to create a strong classifier by adjusting the weights of training samples
- By default it uses the **DecisionStump**



- The default parameters



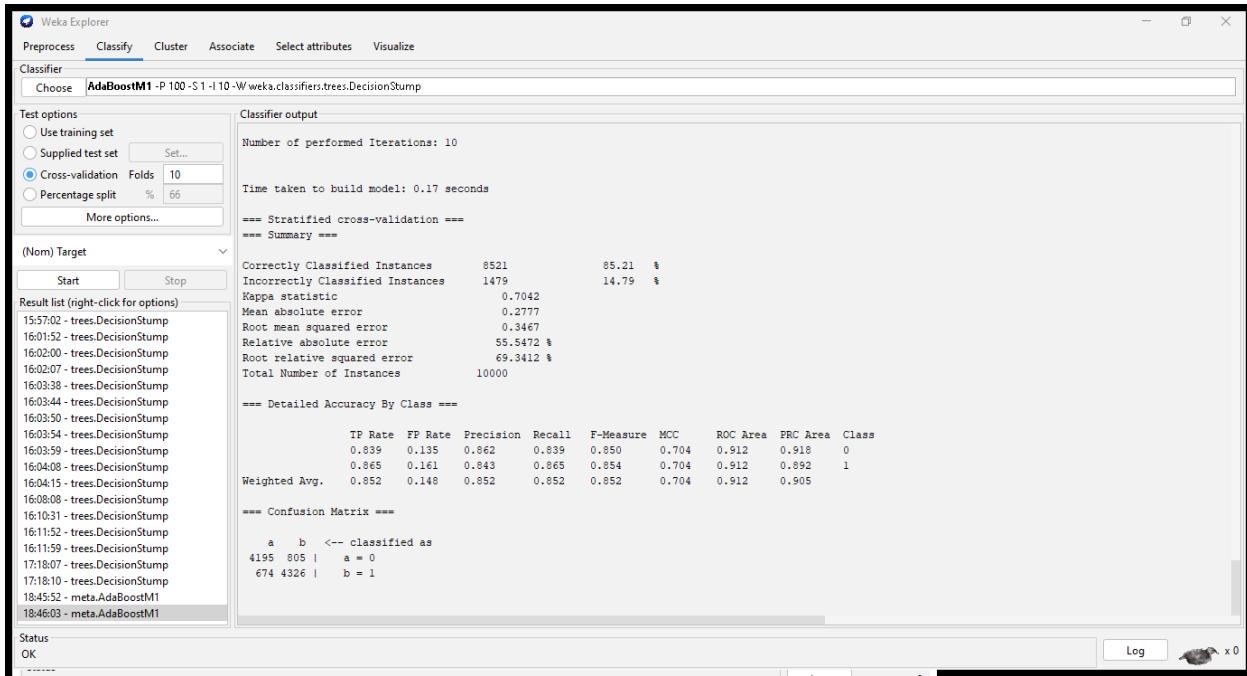
- The result:



The accuracy: 86.38%

The confusion matrix:

	Predicted 0	Predicted 1
Actual 0	1427	279
Actual 1	184	1510

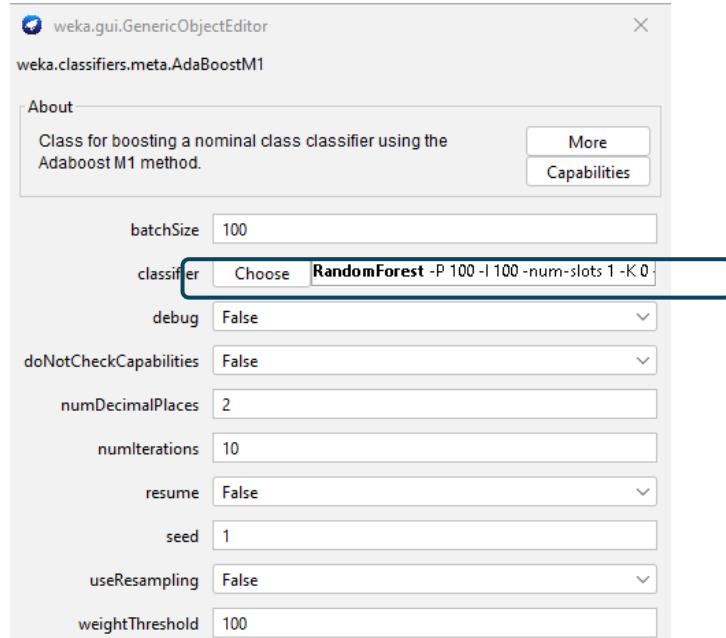


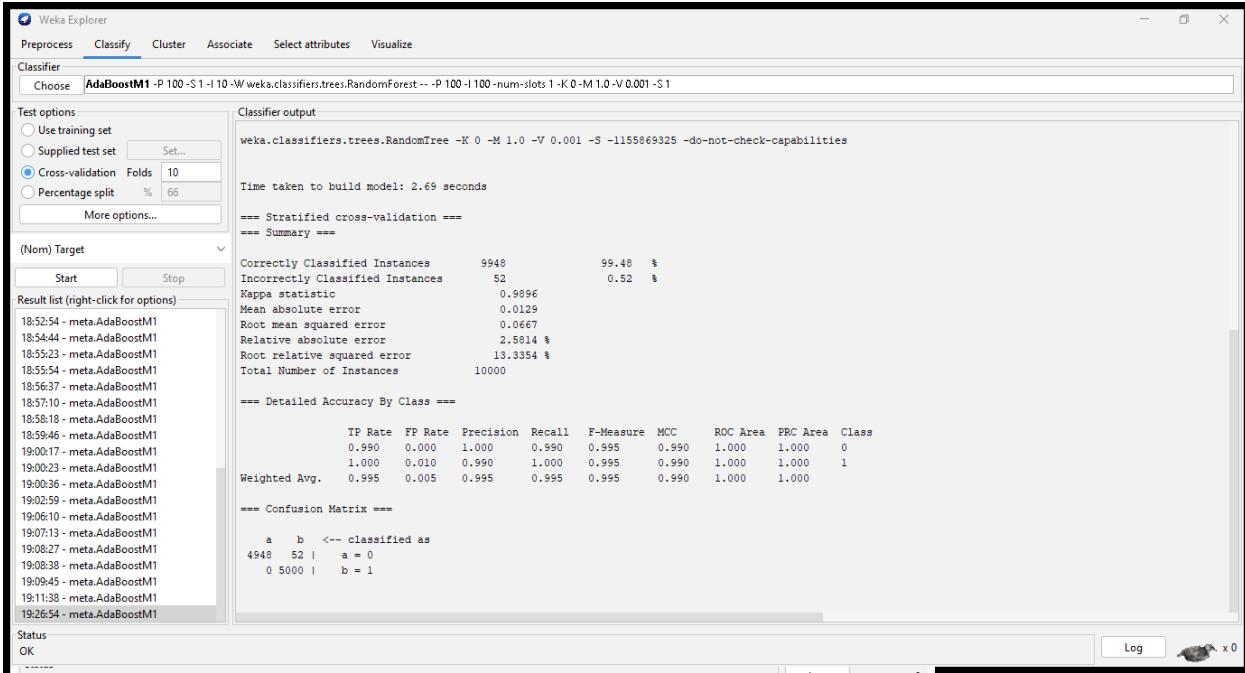
The confusion matrix:

	Predicted 0	Predicted 1
Actual 0	4195	805
Actual 1	674	4326

The algorithm with different parameters:

- When I changed the classifier to **RandomForest**, it gave higher accuracy than the default setting, but it took more time to run.
- **classifier** refers to the base classifier that AdaBoostM1 uses for boosting. It is the model that AdaBoostM1 enhances by adjusting the weights of training samples.

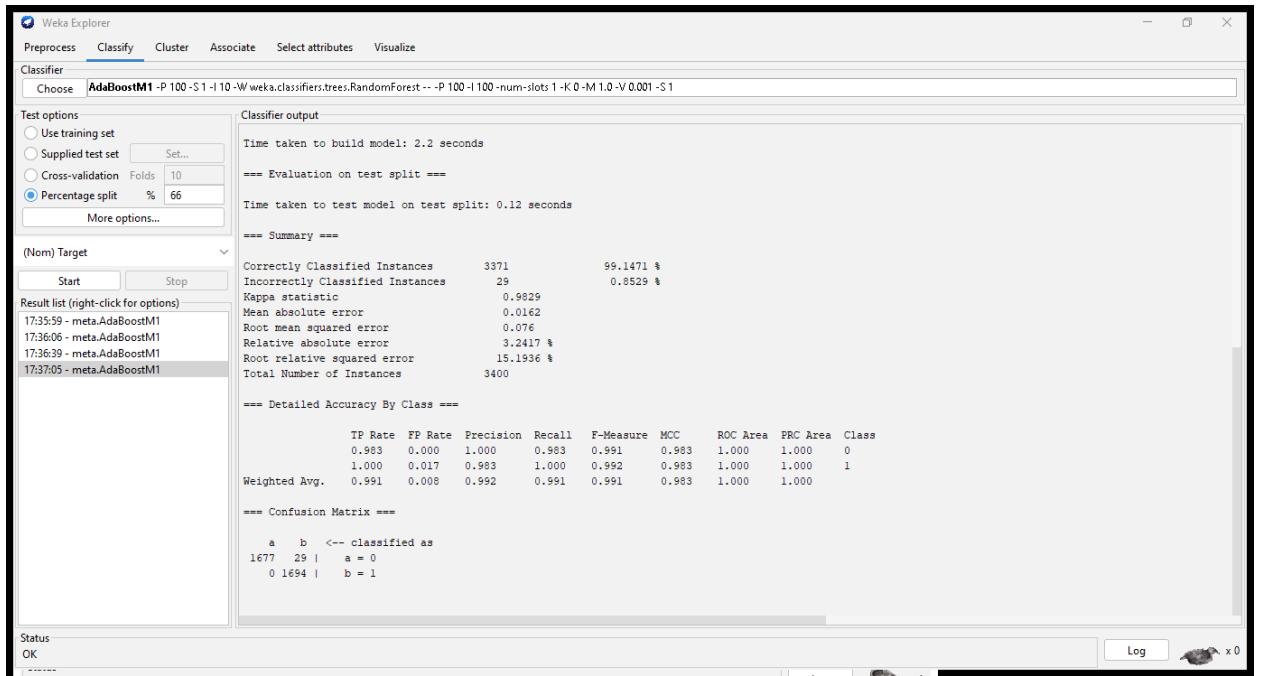




The accuracy: 99.48%

The confusion matrix:

	Predicted 0	Predicted 1
Actual 0	4948	52
Actual 1	0	5000



The accuracy: 99.14%

The confusion matrix:

	Predicted 0	Predicted 1
Actual 0	1677	29
Actual 1	0	1694

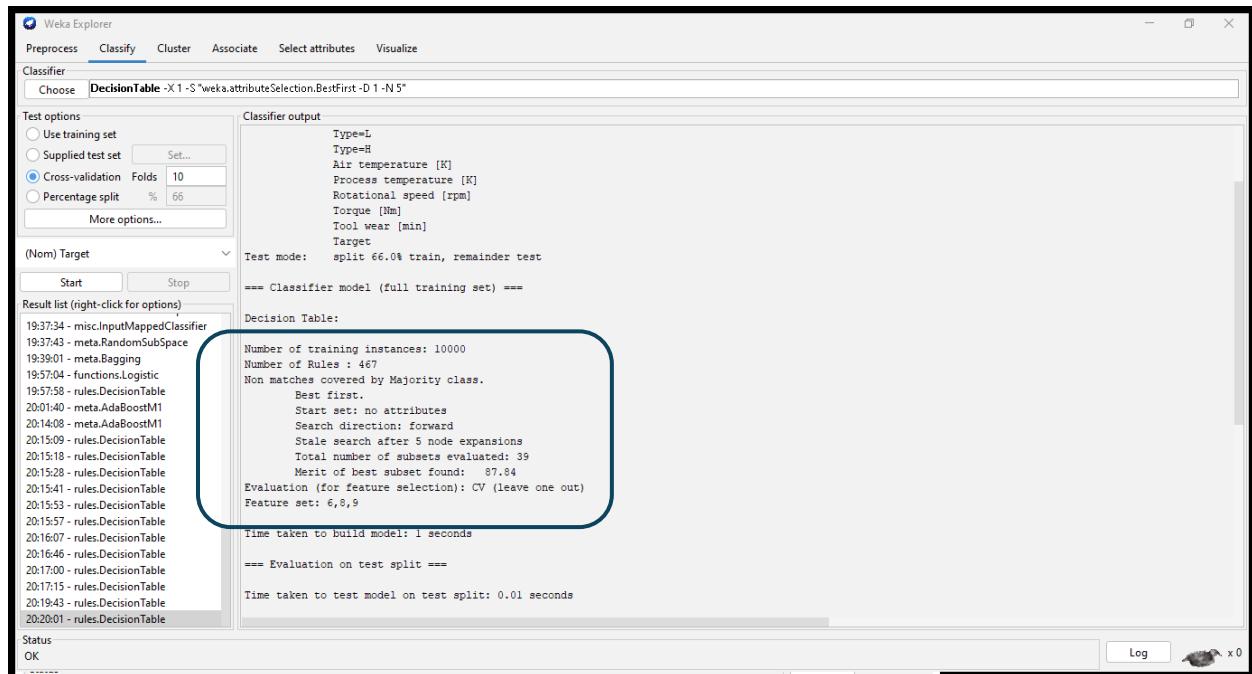
Visualize classifier Error:



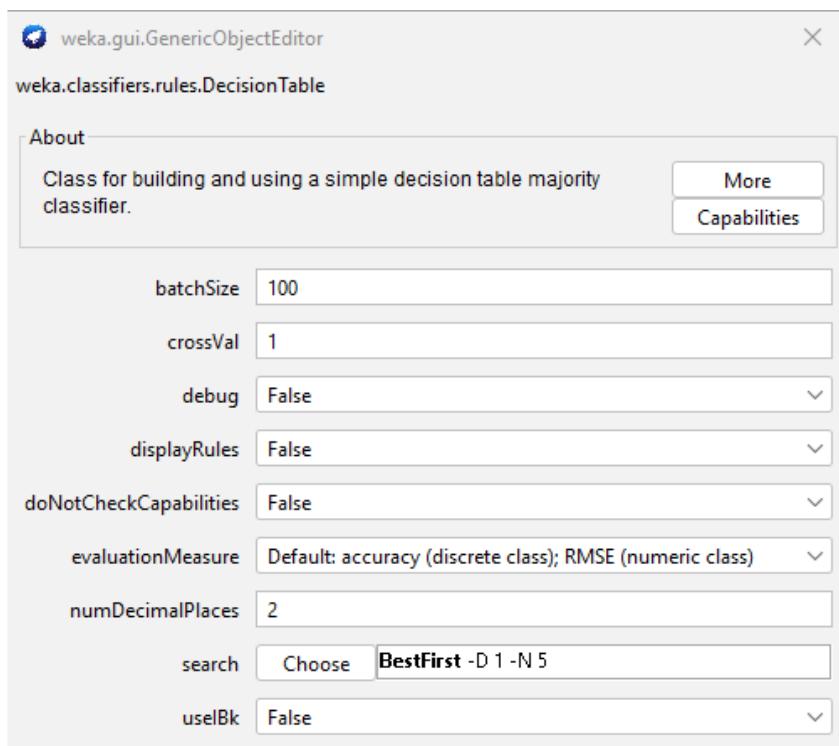
7) DecisionTable:

Class for building and using a simple decision table majority classifier

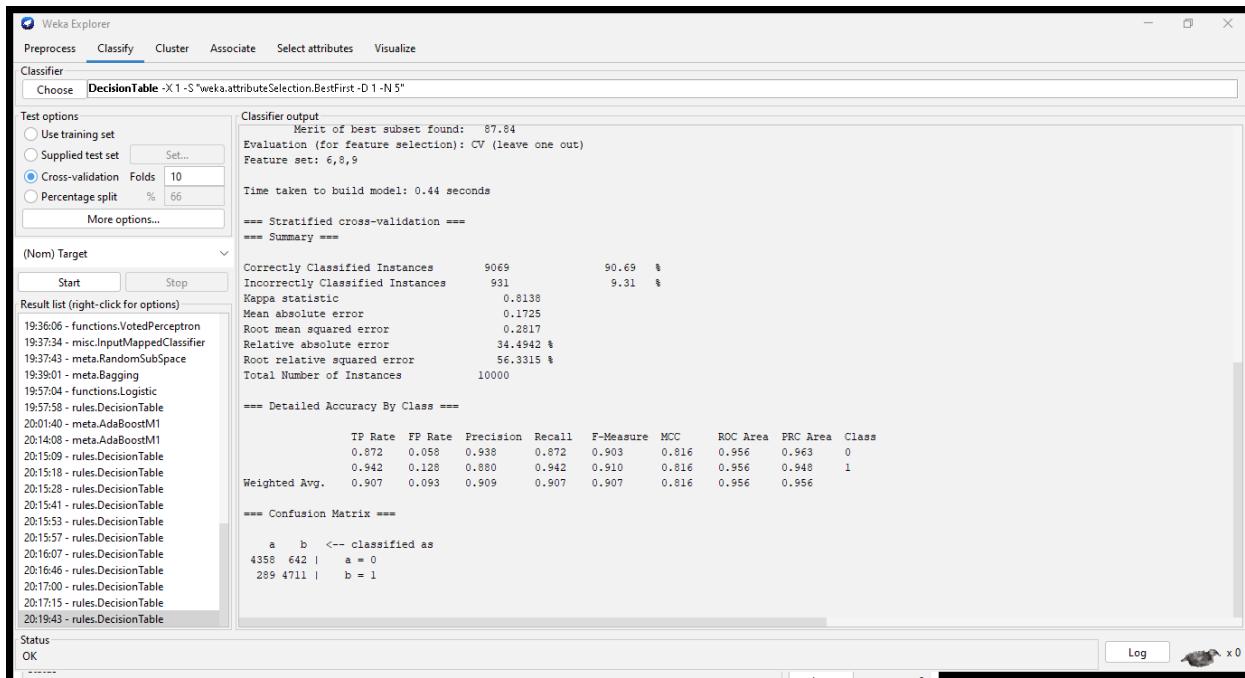
- Total **467 rules** were generated from the training data.
- If no rule matches, it assigns the **majority class**.
- Selected only **3 important features** (features 6, 8, and 9).



- The default parameters:



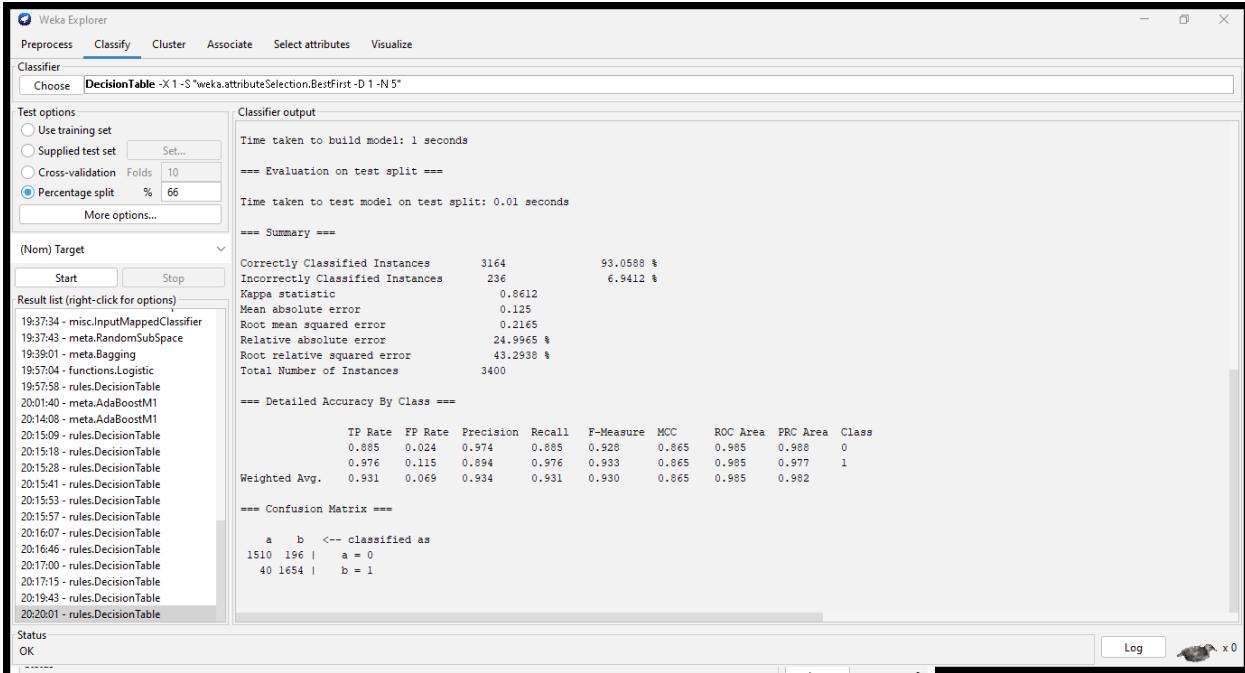
The result:



The accuracy: 90.69%

The confusion matrix:

	Predicted 0	Predicted 1
Actual 0	4358	642
Actual 1	289	4711



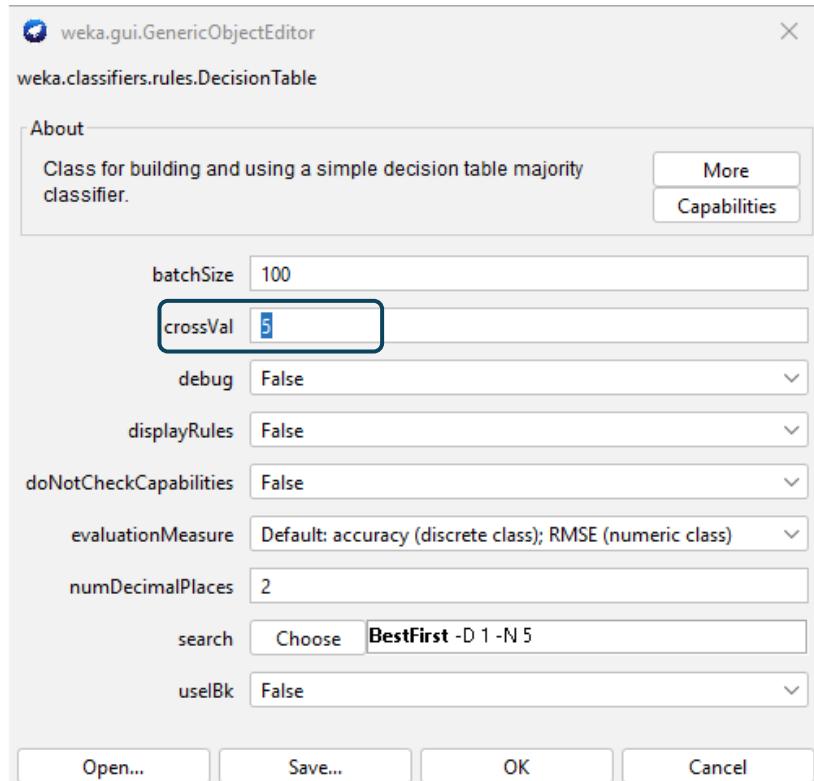
The accuracy: 93.05%

The confusion matrix:

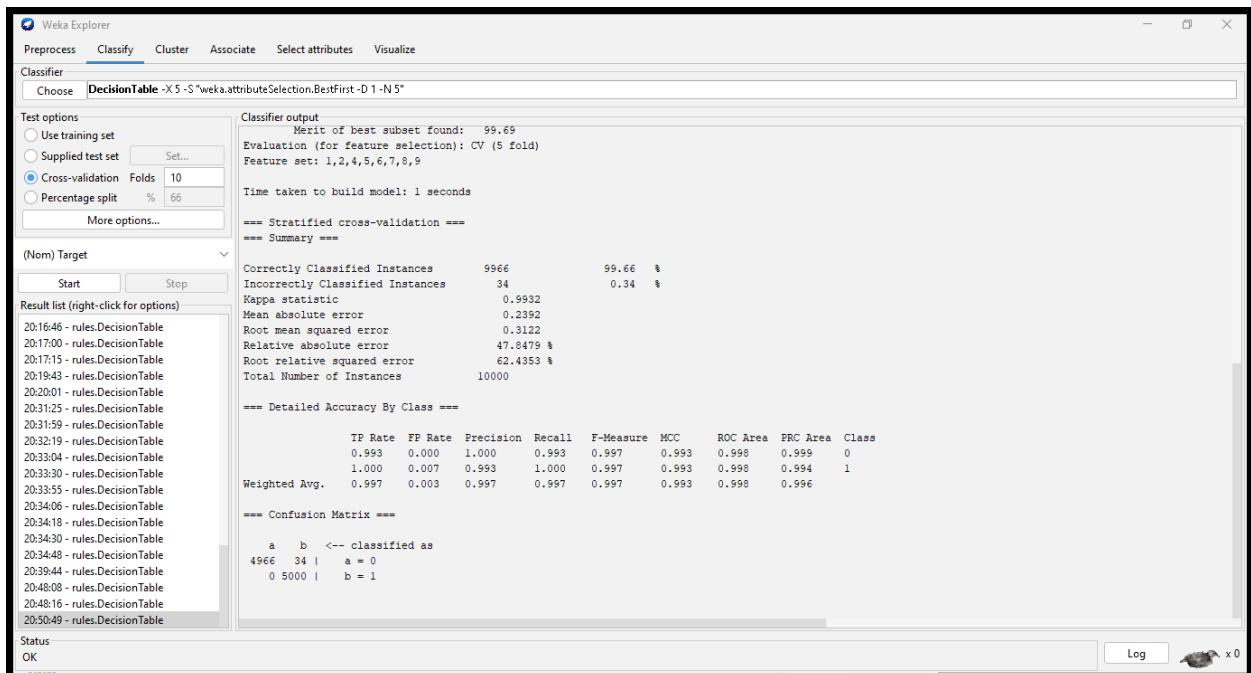
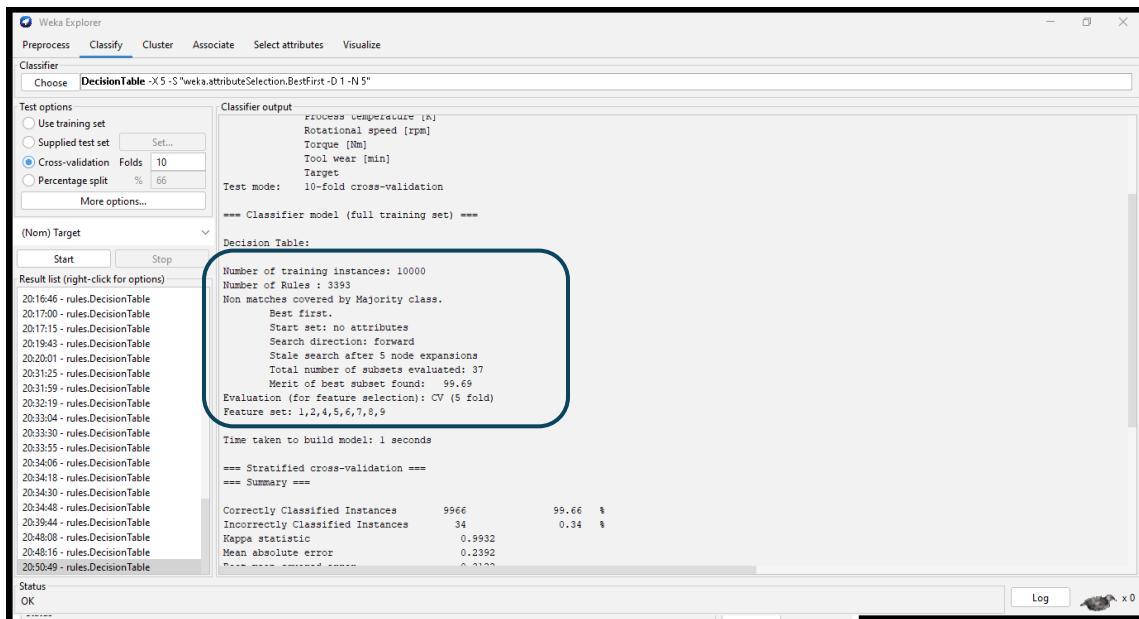
	Predicted 0	Predicted 1
Actual 0	1510	196
Actual 1	40	1654

- Percentage spilt better than cross validation

- **The algorithm with different parameters:**
- When I changed the **crossVal** parameter and increased its value (**from 1 to 5**), the algorithm showed **better** performance compared to the default setting.
- **crossVal** determines how many times cross-validation is used during feature selection. Increasing it allows for a more thorough evaluation of features, which can improve the model's accuracy.



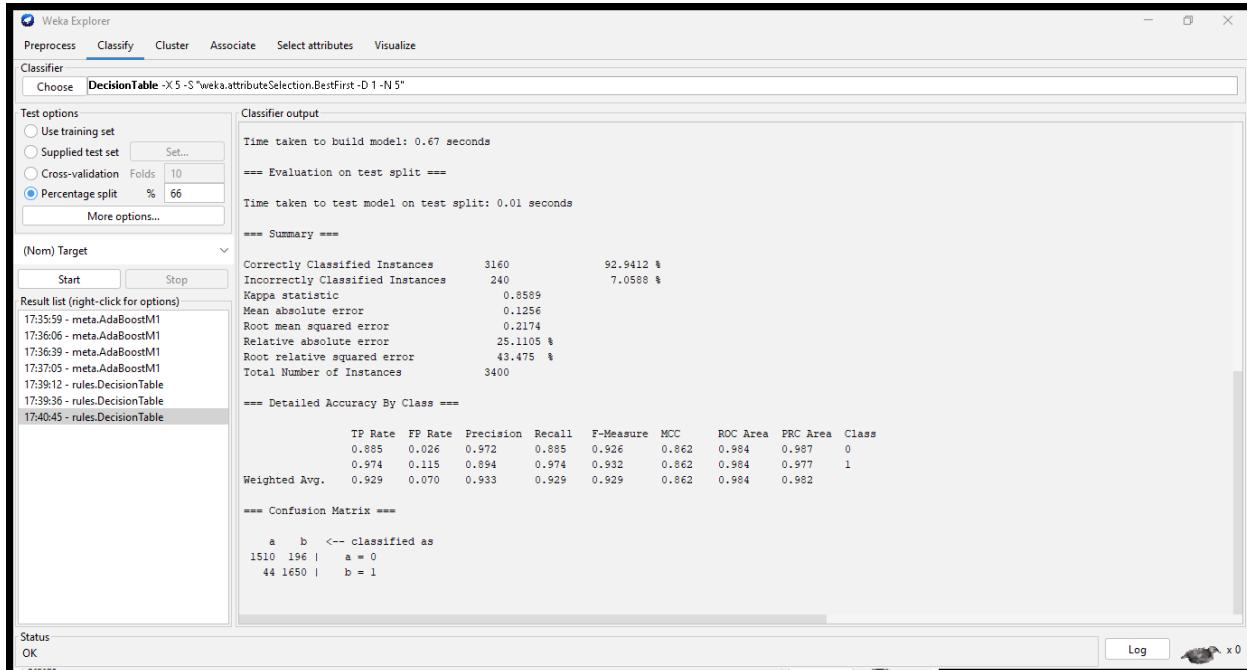
- Total **3393 rules** were generated from the training data.
- If no rule matches, it assigns the **majority class**.
- Selected 8 **important features** (features 1,2,4,5,6,8, and 9).



The accuracy: 99.66%

The confusion matrix:

	Predicted 0	Predicted 1
Actual 0	4966	34
Actual 1	0	5000



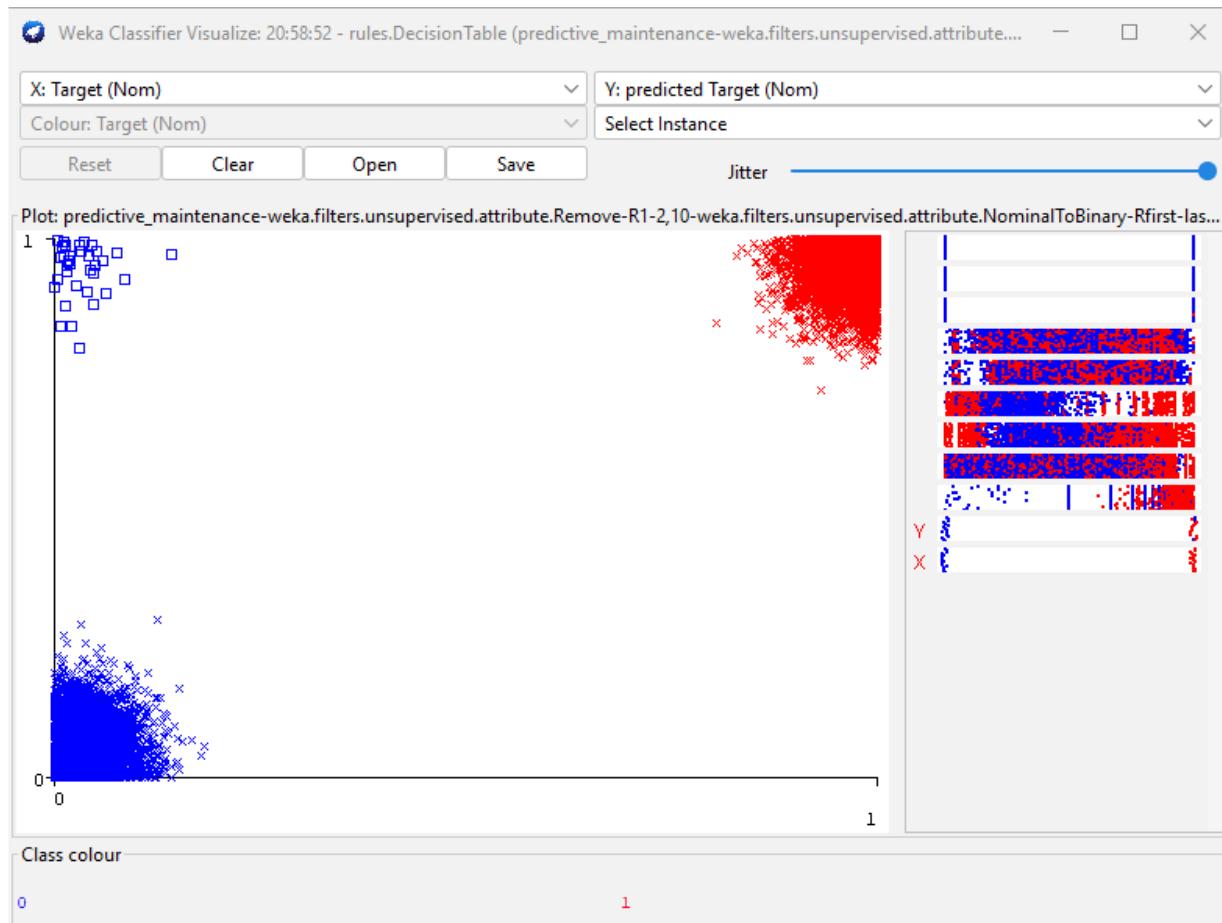
The accuracy: 92.94%

The confusion matrix:

	Predicted 0	Predicted 1
Actual 0	1510	196
Actual 1	44	1650

- At first, with the default parameters, the Percentage Split method gave a higher accuracy than Cross Validation. But after I changed the algorithm's parameters, Cross Validation gave a better result.

- **Visualize classifier Error:**



- **Final result:**

- The **DecisionStump** algorithm achieved an accuracy of **76.97%**.
- The **AdaBoostM1** algorithm achieved an accuracy of **99.48%**.
- The **DecisionTable** algorithm achieved an accuracy of **99.66%**.

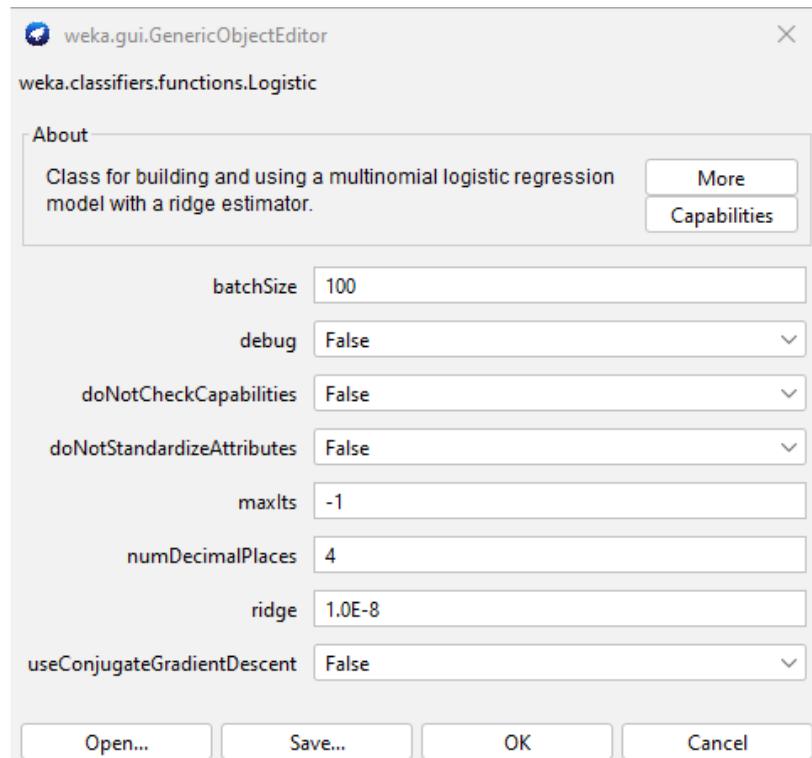
This indicates that the **DecisionTable** algorithm performed the **best** among the three.

8) Logistic Regression (LR)

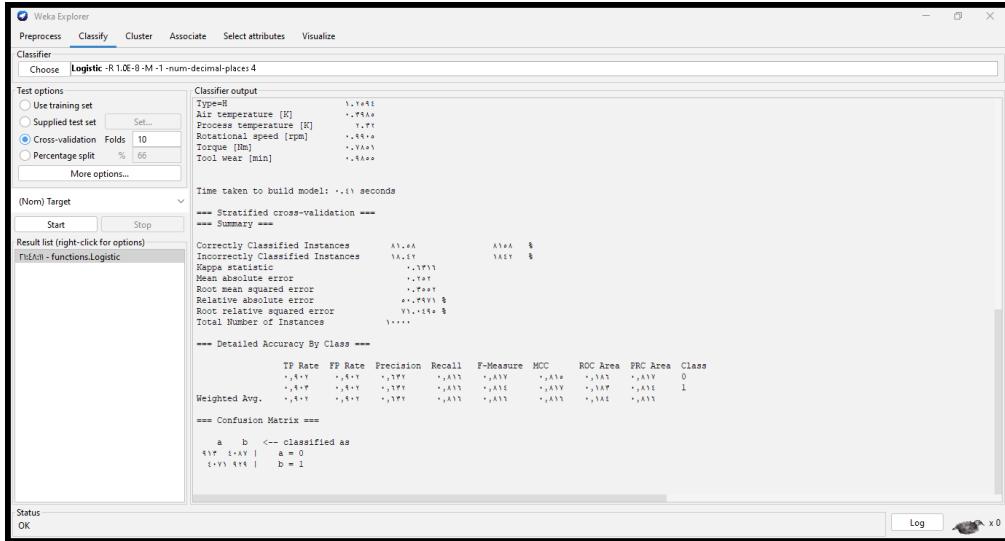
Logistic Regression is a classification method used to predict whether the outcome is 0 or 1 (such as success or failure). It aims to learn the relationship between the input features in order to estimate the probability that a given instance belongs to a particular class or not.

In our data, we applied Logistic Regression to predict whether a machine would fail or succeed based on features like air temperature, rotational speed, and torque.

- The default parameters:



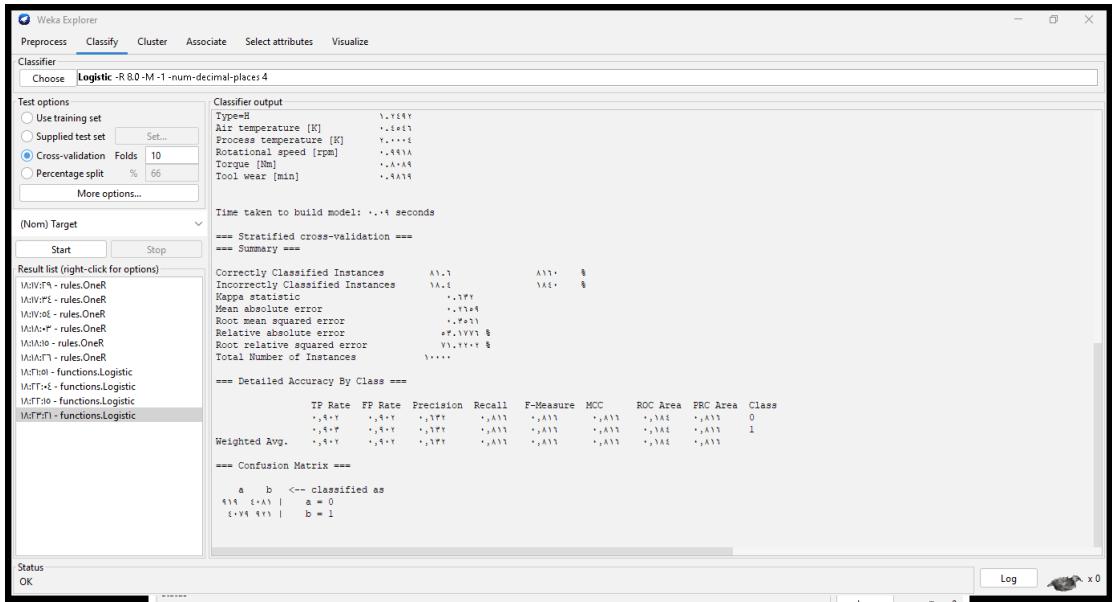
- The result



The accuracy: 81.58%

The confusion matrix:

	Predicted 0	Predicted 1
Actual 0	4087	913
Actual 1	929	4071

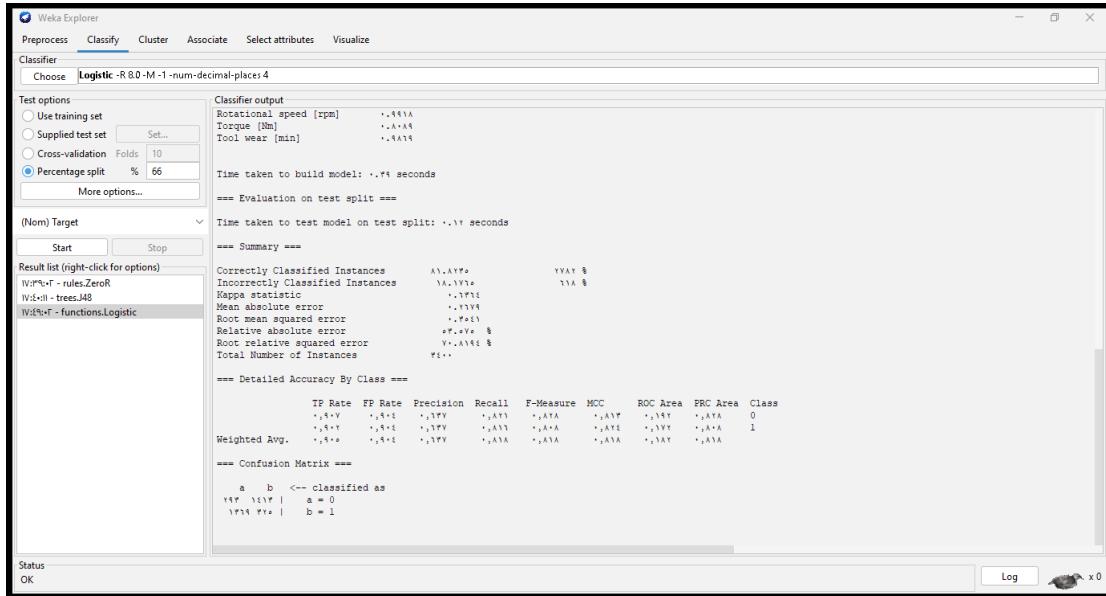


The accuracy : 81.6%

The confusion matrix:

	Predicted 0	Predicted 1
Actual 0	4081	919
Actual 1	921	4079

- Percentage split



The accuracy : 81.8%

The confusion matrix:

	Predicted 0	Predicted 1
Actual 0	1413	293
Actual 1	325	1369

The algorithm with different parameters:

- When I changed the ridge parameter (from 1.0E-8 to 8.0) and increased its value, the algorithm gave better performance compared to the default setting."
- **ridge** Set the Ridge value in the log-likelihood.

batchSize	100
debug	False
doNotCheckCapabilities	False
doNotStandardizeAttributes	False
maxItls	-1
numDecimalPlaces	4
ridge	8.0
useConjugateGradientDescent	False

- Visualize classifier Error:



9) (J48) Tree :

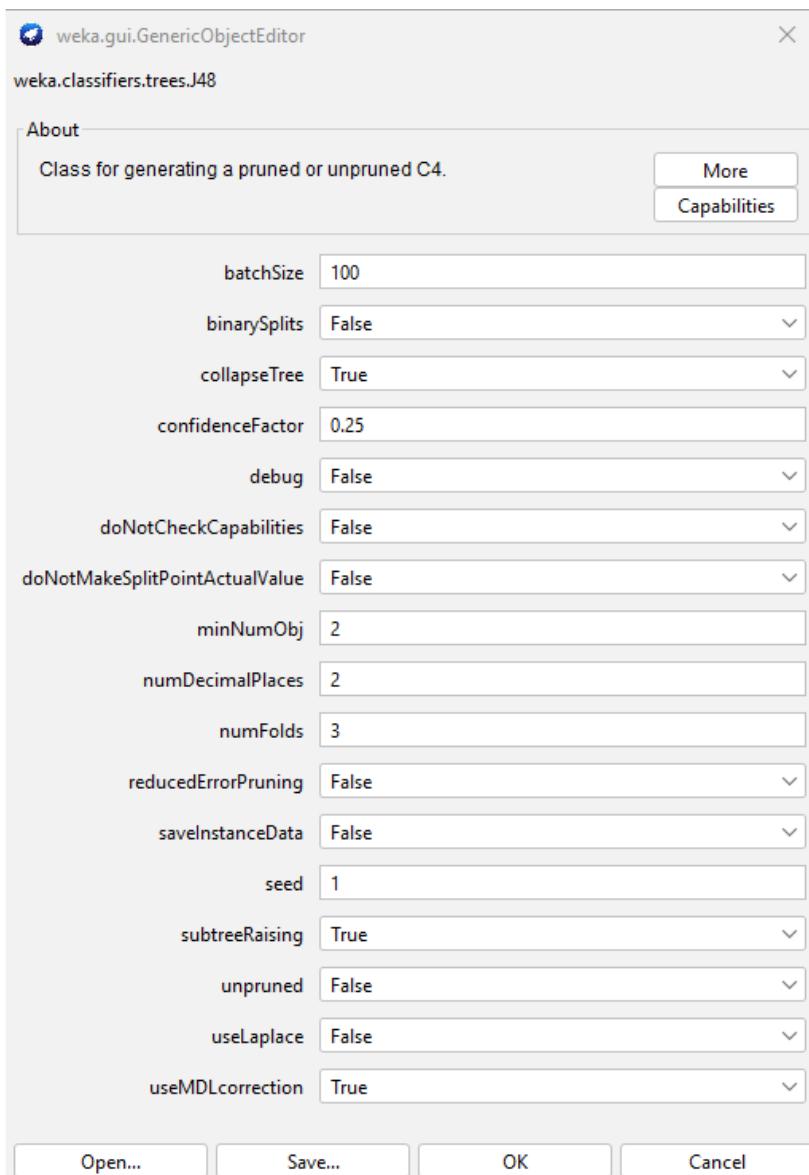
J48 classifier is a Decision tree algorithm that builds a decision tree based on the attribute values of the dataset.

The J48 classifier is popular because it produces interpretable models and can handle both numeric and nominal attributes.

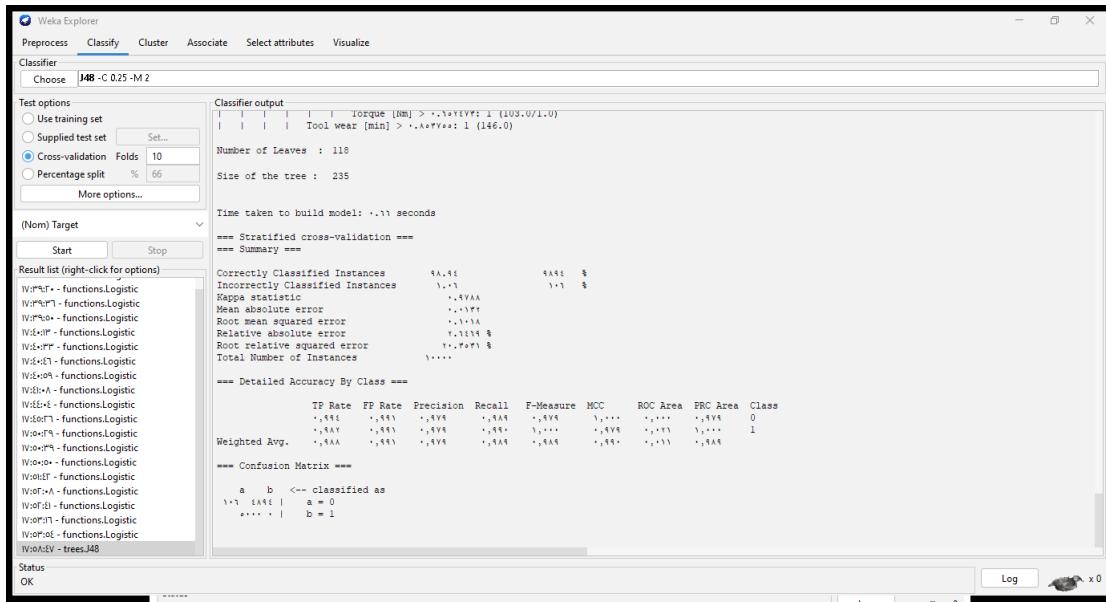
It supports pruning techniques to reduce overfitting and provides options for parameter tuning to optimize the model's accuracy.

In our data , we used J48 to predict whether the machine would fail or work.

- The default parameters:



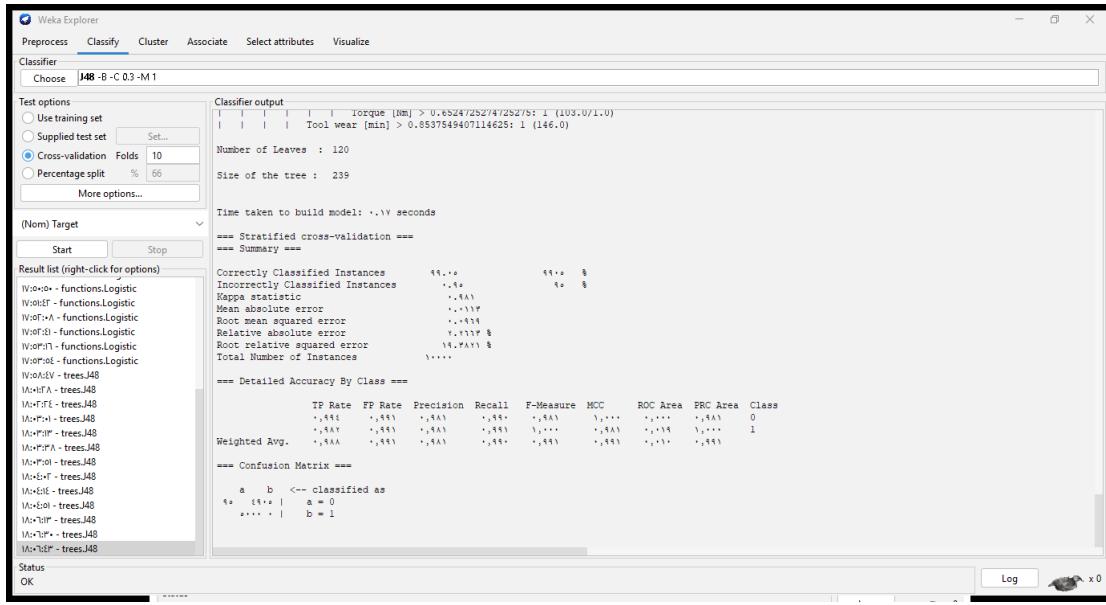
- The result



The accuracy: 98.94%

The confusion matrix:

	Predicted 0	Predicted 1
Actual 0	4894	106
Actual 1	0	5000

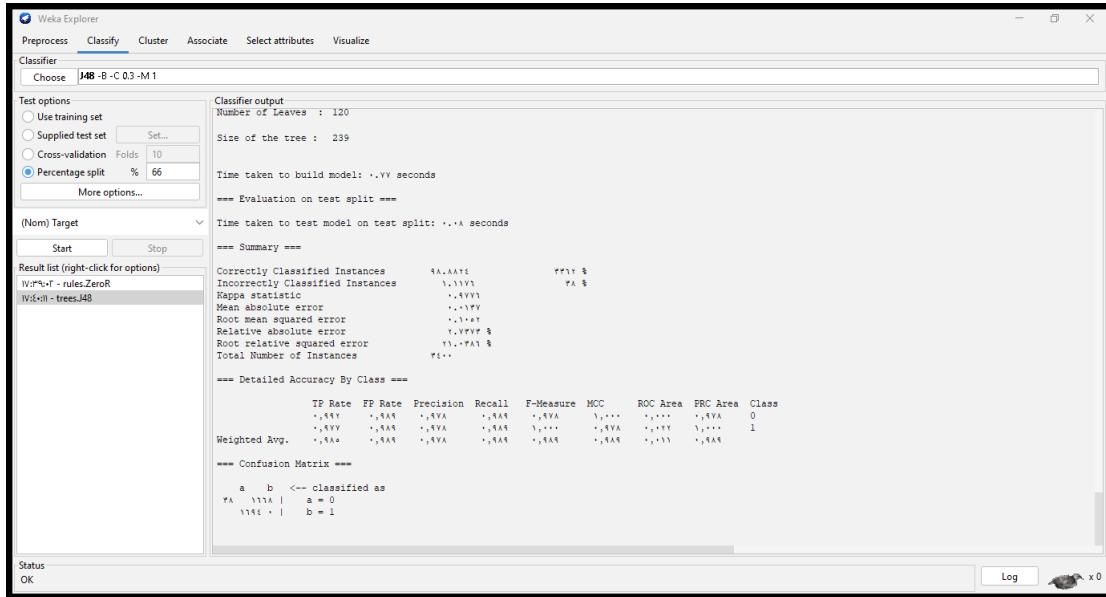


The accuracy : 99.5%

The confusion matrix:

	Predicted 0	Predicted 1
Actual 0	4905	95
Actual 1	0	5000

- Percentage split

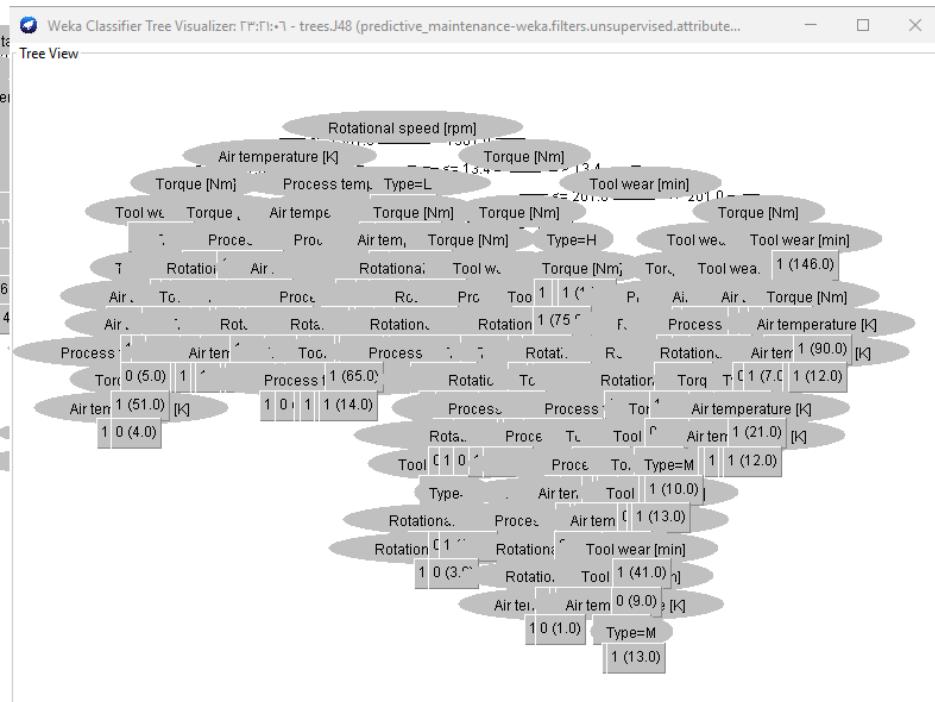


The accuracy : 98.88%

The confusion matrix:

	Predicted 0	Predicted 1
Actual 0	1668	38
Actual 1	0	1694

- Tree J48

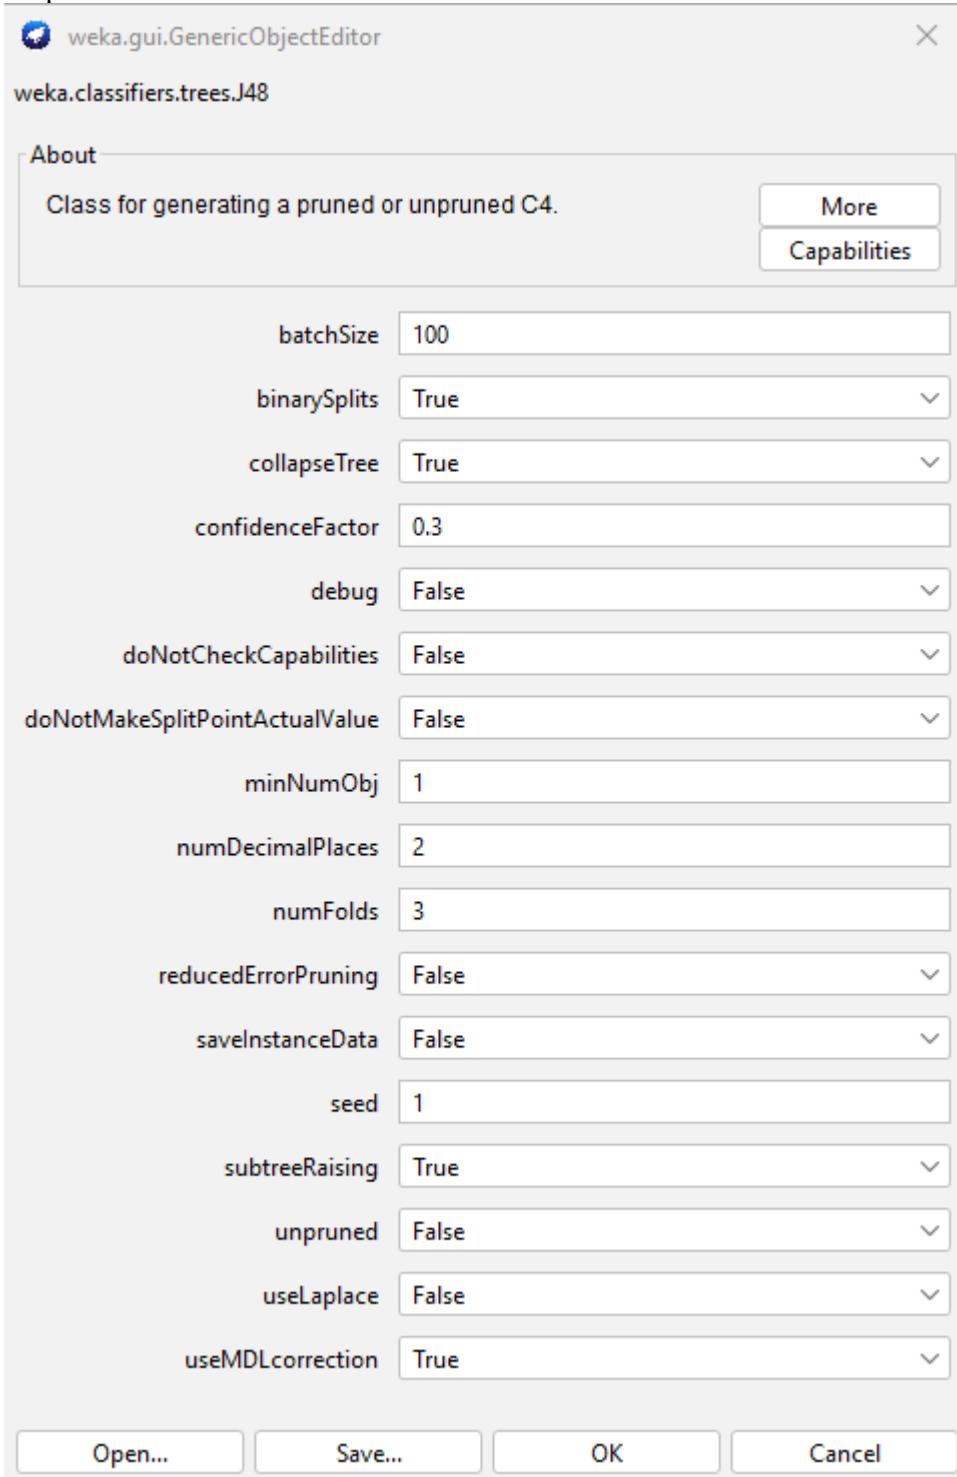


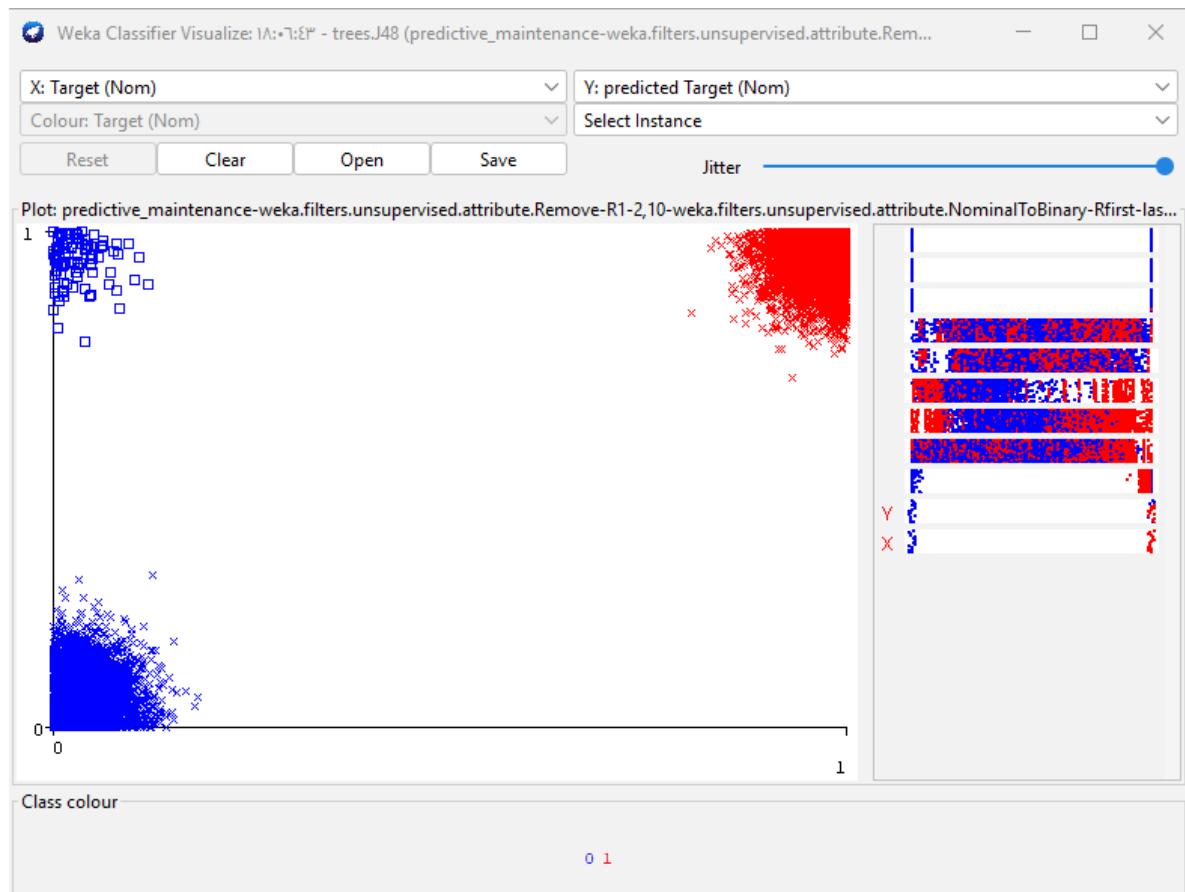
The algorithm with different parameters:

When I changed the binarySplits parameter (from false to true), the confidenceFactor (from 0.25 to 0.3), and the minNumObj (from 2 to 1), the algorithm improved performance."

- **binarySplits**: Whether to use binary splits on nominal attributes when building the trees.
 - **confidenceFactor**: The confidence factor used for pruning (smaller values incur more pruning).
 - **minNumObj**: The minimum number of instances per leaf.

The parameter



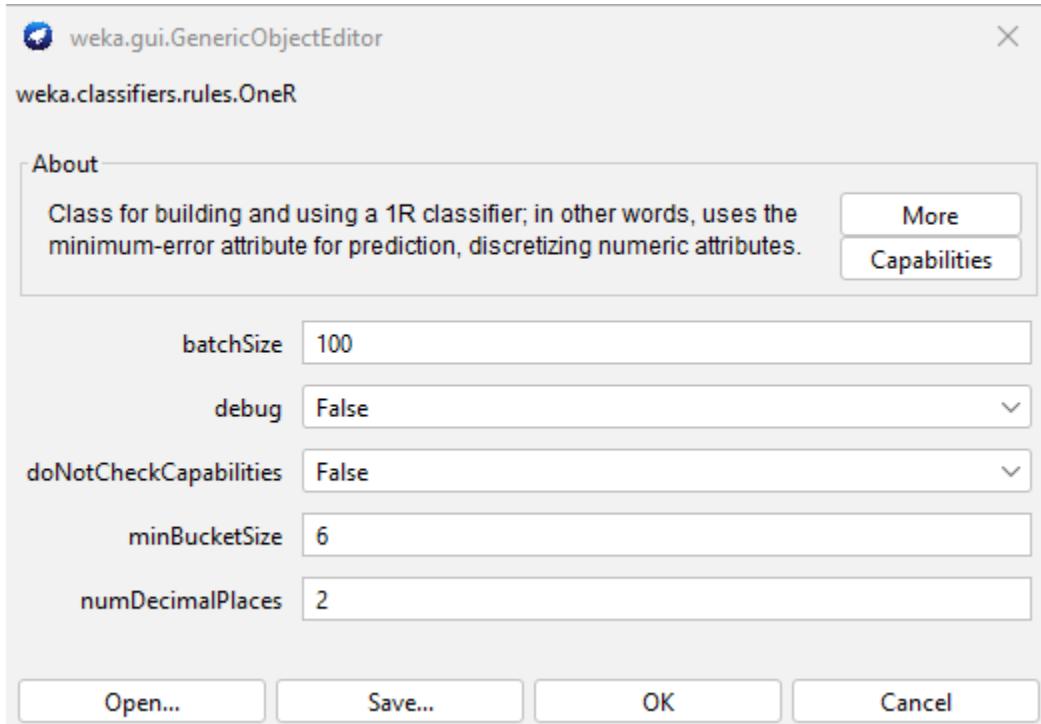


10) OneR:

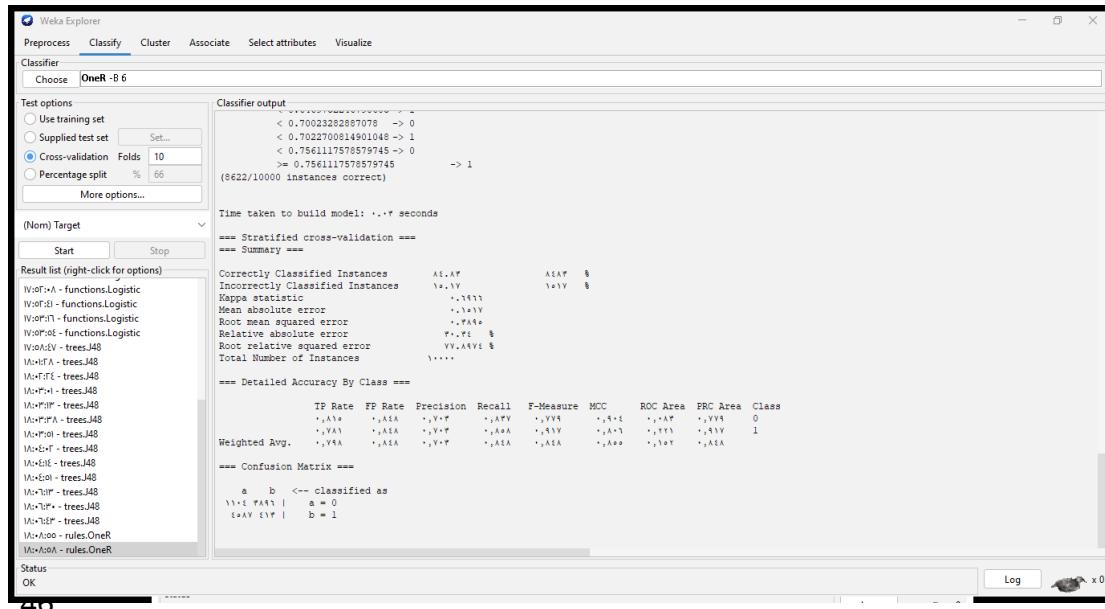
OneR is a simple classification algorithm that learns rules from a single attribute. It is characterized by its simplicity and ease of interpretation. The algorithm examines all attributes and selects the one that results in the lowest classification error.

In our data, rotational speed was chosen as the most important attribute

- The default parameters:



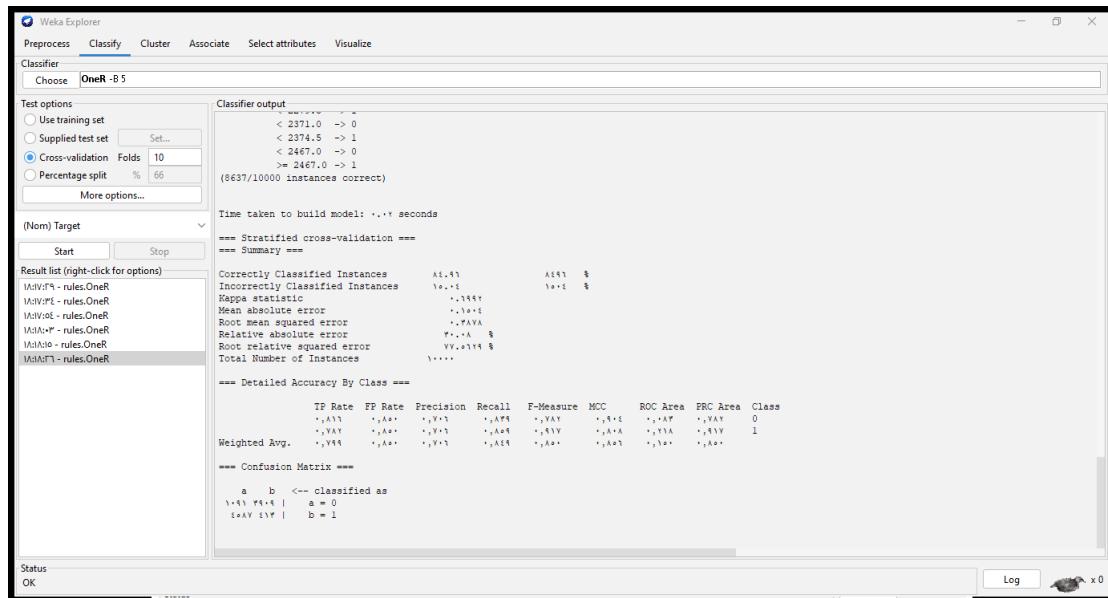
- The result



The accuracy : 84.83%

The confusion matrix:

	Predicted 0	Predicted 1
Actual 0	3891	1104
Actual 1	413	4587

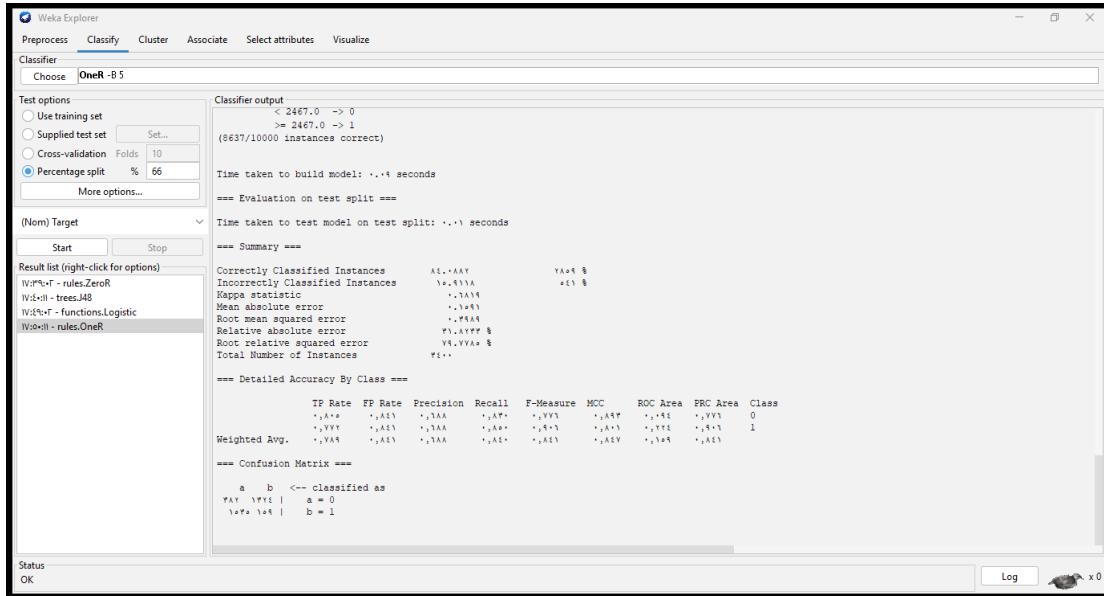


The accuracy : 85.96%

The confusion matrix:

	Predicted 0	Predicted 1
Actual 0	3909	1091
Actual 1	413	4587

- Percentage split



The accuracy : 84.08%

The confusion matrix:

	Predicted 0	Predicted 1
Actual 0	1324	382
Actual 1	159	1535

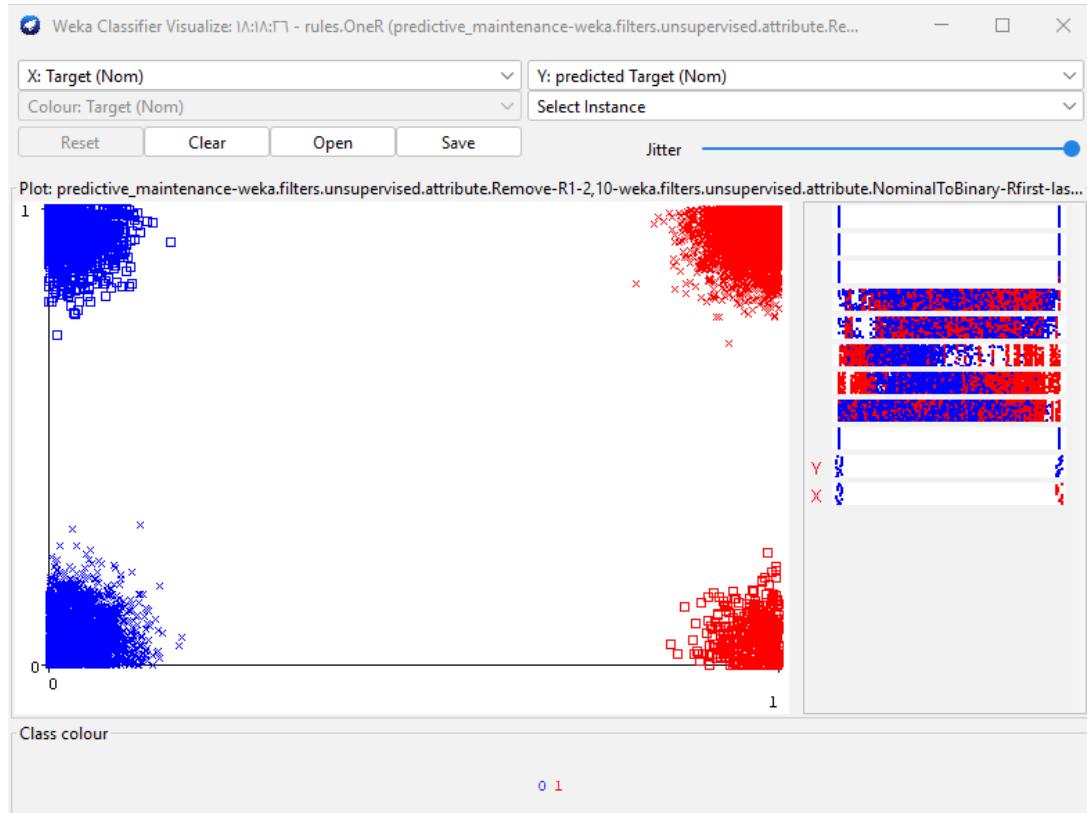
The algorithm with different parameters:

When I changed the minBucketSize parameter (from 6 to 5), the algorithm showed improved performance."

- **minBucketSize**: The minimum bucket size used for discretizing numeric attributes.

batchSize	100
debug	False
doNotCheckCapabilities	False
minBucketSize	5
numDecimalPlaces	2

- Visualize classifier Error:



- **Final result:**

- The Logistic Regression (LR) algorithm achieved an accuracy of 81.6%.
- The Decision Tree (J48) algorithm achieved an accuracy of 99.5%.
- The OneR algorithm achieved an accuracy of 85.96%.

This indicates that the Decision Tree (J48) algorithm performed the best among the three.

11) Simple logistic

The model Focus on all features. It's a linear model that assigns weights to each feature to predict probabilities

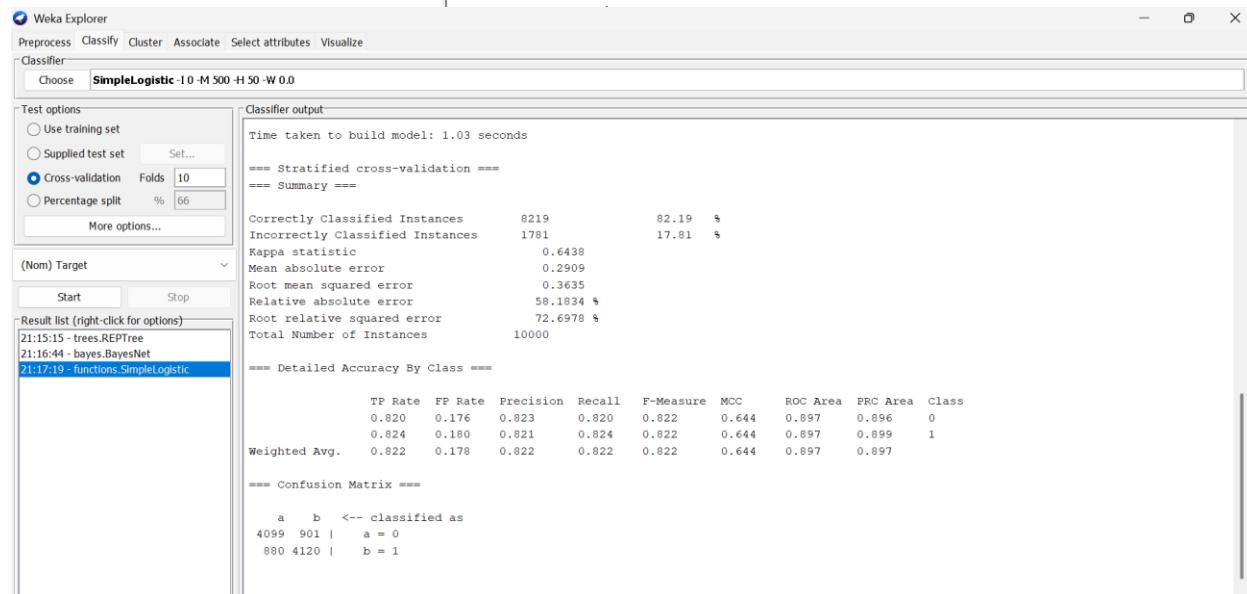
Focus on our data: Likely emphasizes Process temperature, Torque, and Tool wear

Most important features:

- Air temperature
- Torque
- Type L

- The result:

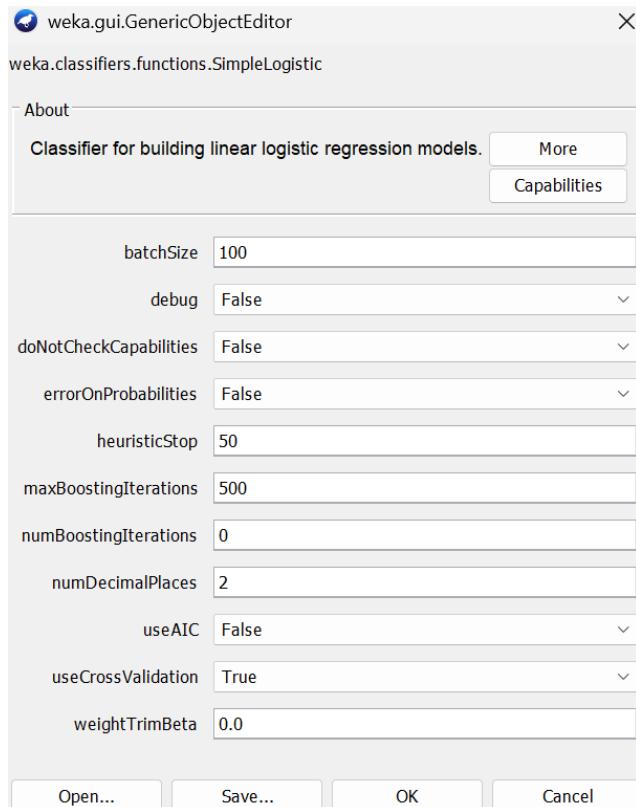
```
SimpleLogistic:  
  
Class 0 :  
39.45 +  
[Type=L] * -0.16 +  
[Air temperature [K]] * -0.24 +  
[Process temperature [K]] * 0.14 +  
[Rotational speed [rpm]] * -0 +  
[Torque [Nm]] * -0.09 +  
[Tool wear [min]] * -0  
  
Class 1 :  
-39.45 +  
[Type=L] * 0.16 +  
[Air temperature [K]] * 0.24 +  
[Process temperature [K]] * -0.14 +  
[Rotational speed [rpm]] * 0 +  
[Torque [Nm]] * 0.09 +
```



The confusion matrix:

	Predicted 0	Predicted 1
Actual 0	4099	901
Actual 1	880	4120

- Default parameter



- Different parameters

When trying to change the heuristicsStop so the model could try to learn and optimize more than 50 iterations the accuracy did not change, and the same thing happened to maxBoostingIterations when increasing the number of iterations, the accuracy has stayed to 82.19%.

- Visualize errors



12) Bayes Network

The model can use all features but only models dependencies between them. Some features may be conditionally independent and ignored.

On our data every feature is directly linked to Target — so all are used.

```
Bayes Network Classifier
not using ADTree
#attributes=9 #classindex=8
Network structure (nodes followed by parents)
Type=M(2): Target
Type=L(2): Target
Type=H(2): Target
Air temperature [K](5): Target
Process temperature [K](9): Target
Rotational speed [rpm](29): Target
Torque [Nm](35): Target
Tool wear [min](10): Target
Target(2):
LogScore Bayes: -102153.74379441359
LogScore BDeu: -102841.32867280956
LogScore MDL: -102850.64282791069
LogScore ENTROPY: -102053.94838573477
LogScore AIC: -102226.94838573477
```

```
Time taken to build model: 0.22 seconds
```

- The result:

```
Correctly Classified Instances      8941          89.41 %
Incorrectly Classified Instances   1059          10.59 %
Kappa statistic                   0.7882
Mean absolute error               0.129
Root mean squared error           0.2796
Relative absolute error            25.81 %
Root relative squared error      55.914 %
Total Number of Instances        10000

==== Detailed Accuracy By Class ====

      TP Rate  FP Rate  Precision  Recall   F-Measure  MCC     ROC Area  PRC Area  Class
          0.860    0.072    0.923    0.860    0.890    0.790    0.963    0.966    0
          0.928    0.140    0.869    0.928    0.898    0.790    0.963    0.960    1
Weighted Avg.      0.894    0.106    0.896    0.894    0.894    0.790    0.963    0.963

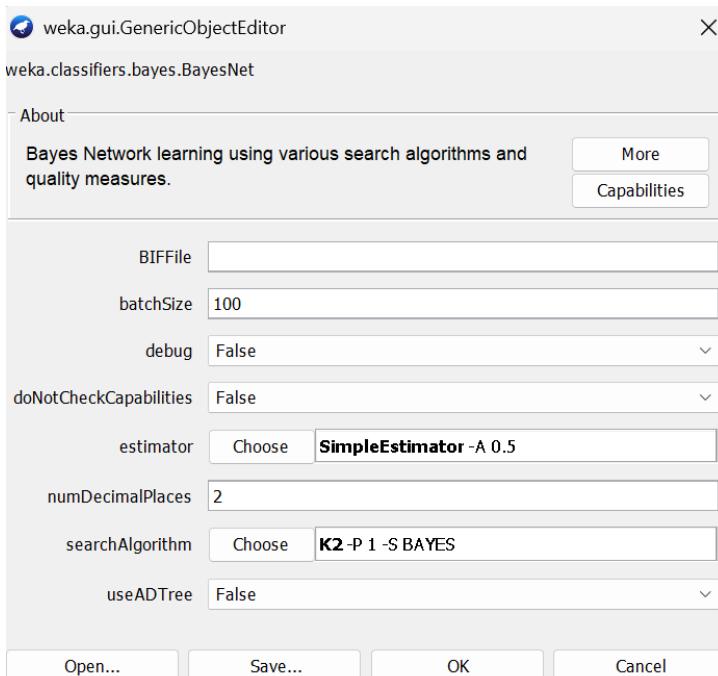
==== Confusion Matrix ====

      a      b  <-- classified as
4300  700 |  a = 0
359 4641 |  b = 1
```

The confusion matrix:

	Predicted 0	Predicted 1
Actual 0	4300	700
Actual 1	359	4641

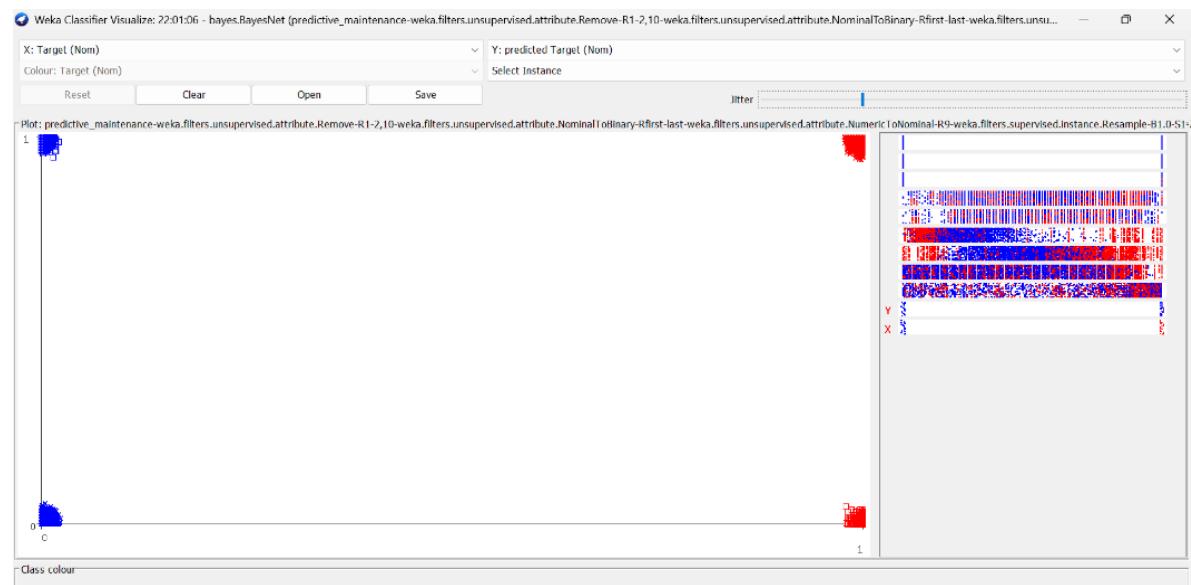
-Default parameters



-Different parameters

When decreasing the Estimator number to 0.2 and to 0.1 the accuracy has **increase to 89.46%**.

-Visualize errors

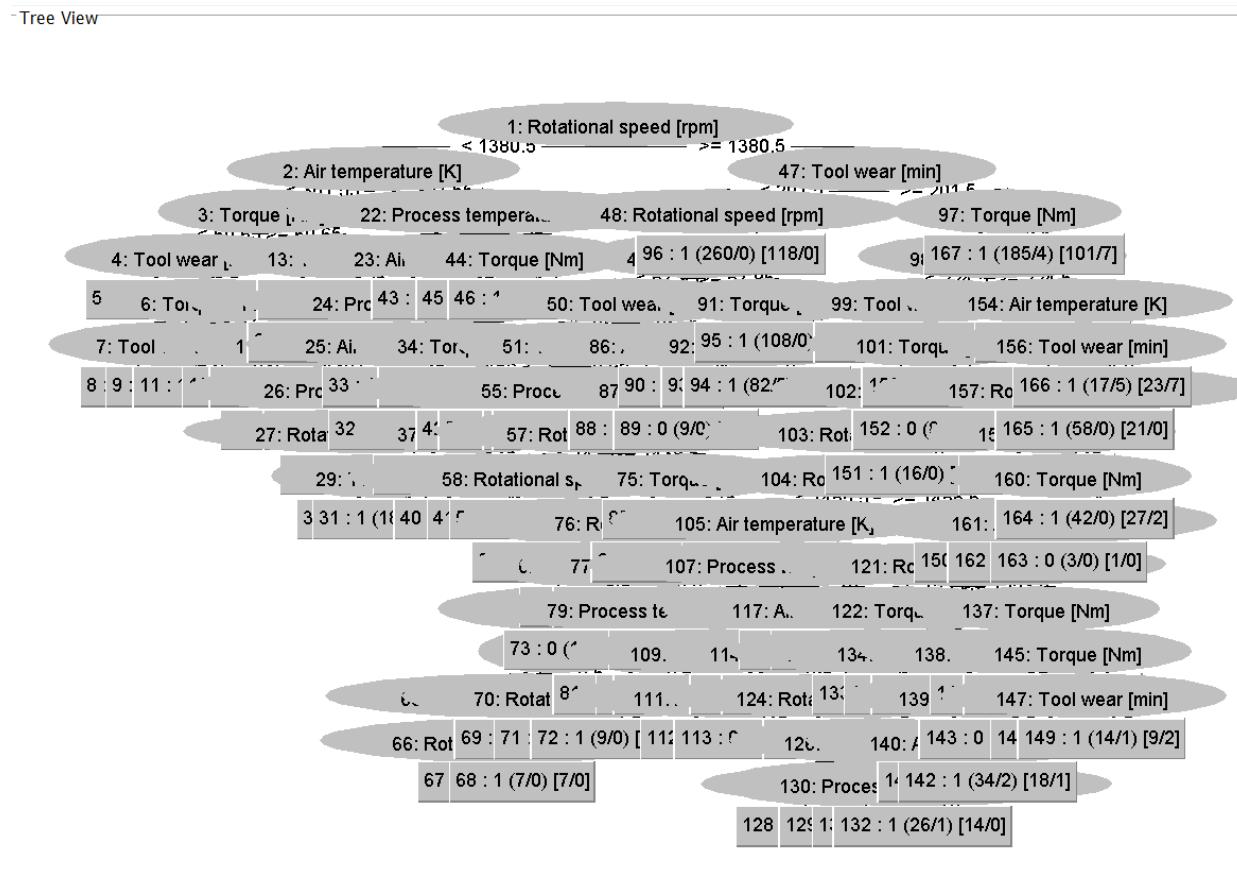


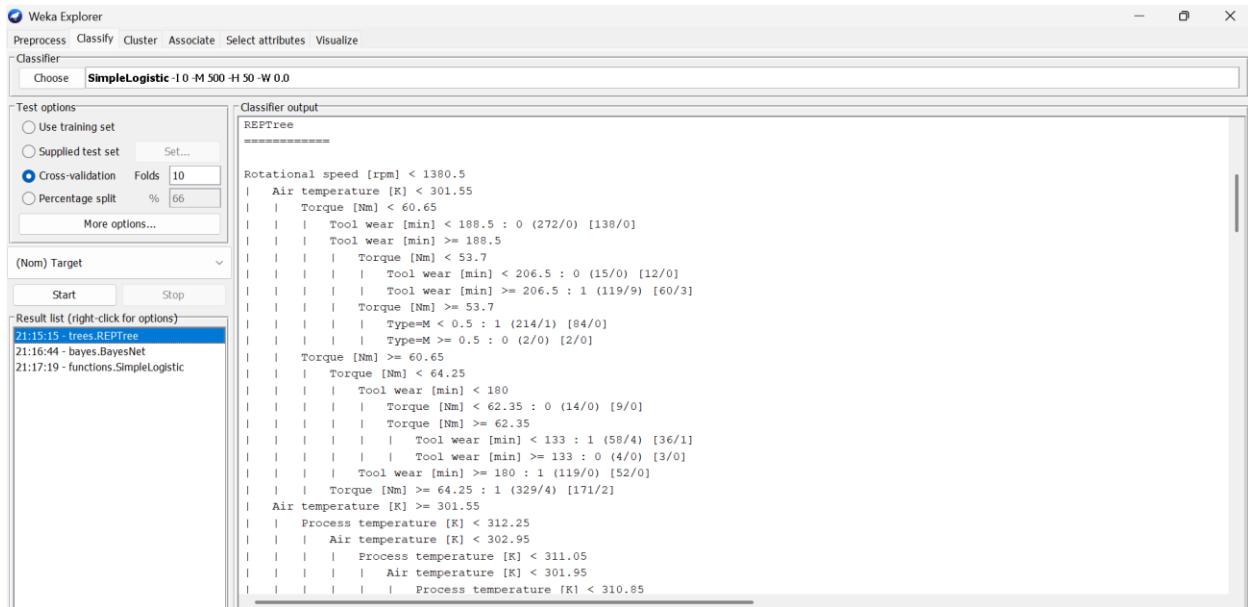
13) REP tree

This model uses a subset of features, based on decision splits that most reduce error.

Focus on our data: Features used heavily:

- Rotational speed and Tool wear are top decision nodes.
 - Torque and Process temperature are used deeper in the tree.
 - Air temperature and Type are used occasionally, less dominant.
 - Focus: Primarily on Rotational speed, Tool wear, Torque, and Process temperature. These features likely provide the highest information gain.





- The result:

Correctly Classified Instances	9832	98.32 %
Incorrectly Classified Instances	168	1.68 %
Kappa statistic	0.9664	
Mean absolute error	0.0247	
Root mean squared error	0.1246	
Relative absolute error	4.9345 %	
Root relative squared error	24.9176 %	
Total Number of Instances	10000	

==== Detailed Accuracy By Class ====

TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
0.968	0.001	0.999	0.968	0.983	0.967	0.992	0.994	0
0.999	0.032	0.969	0.999	0.983	0.967	0.992	0.985	1
Weighted Avg.	0.983	0.017	0.984	0.983	0.983	0.967	0.992	0.989

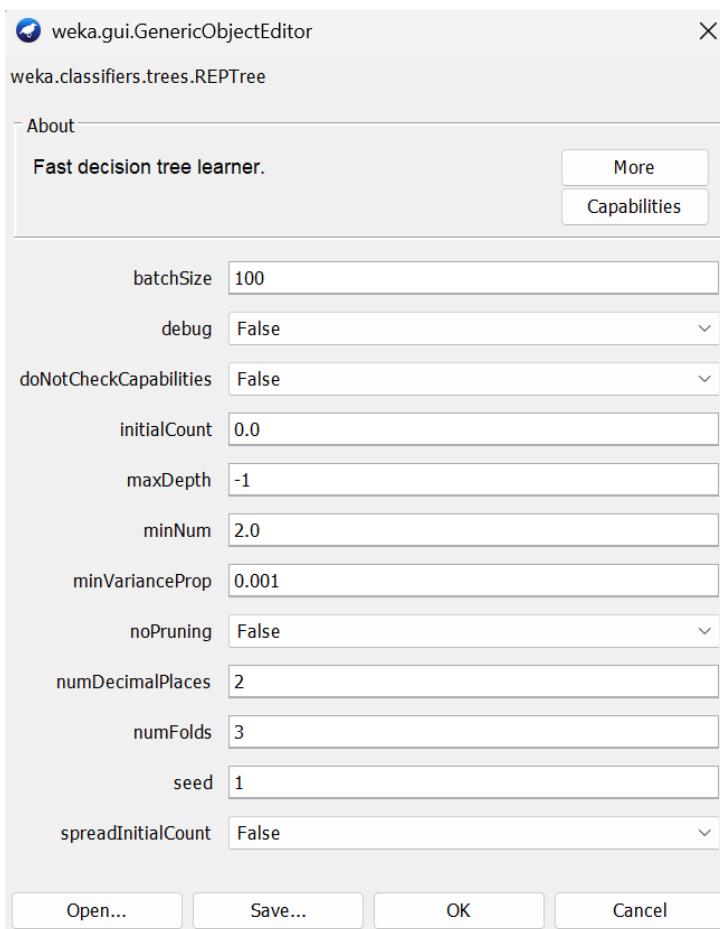
==== Confusion Matrix ====

a	b	<-- classified as
4839	161	a = 0
7	4993	b = 1

The confusion matrix:

	Predicted 0	Predicted 1
Actual 0	4839	161
Actual 1	7	4993

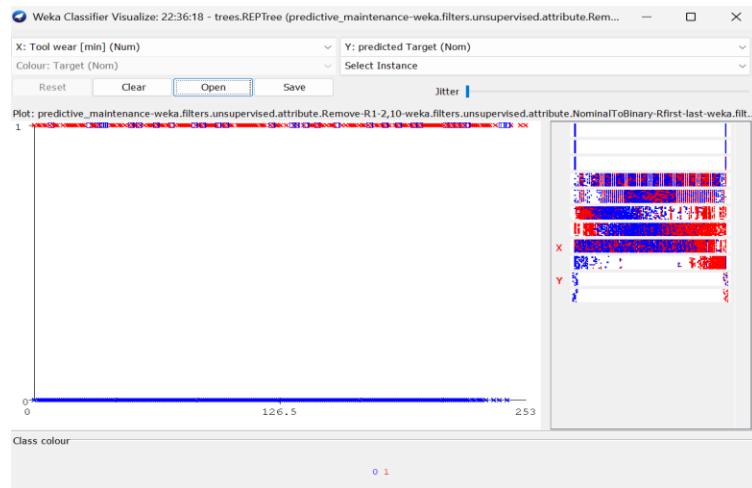
-Default parameter



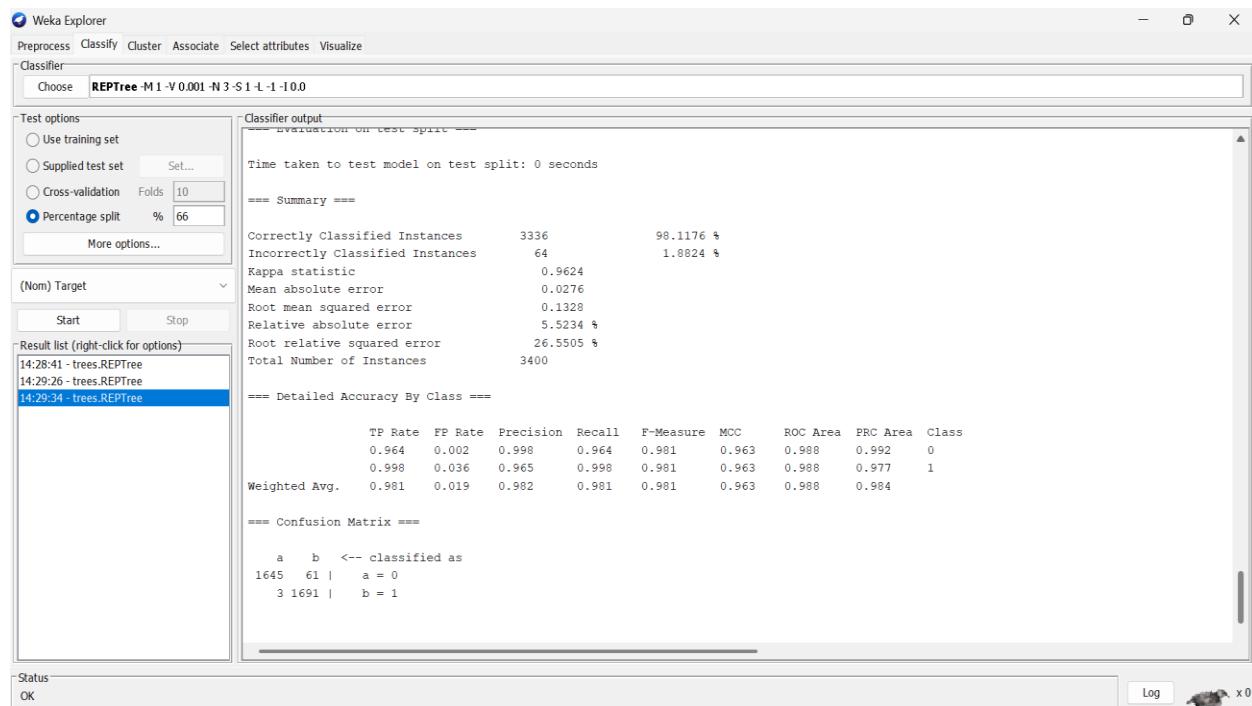
-Different parameter

When changing the minNum to 1, the data becomes more detailed, and the accuracy has increased to 98.51

-Visualize errors



Using cross validation with the best Algorithm accuracy REP Tree 98.51% comparing to Bayes Network and simple logistics algorithms



14) Random Forest:

Random Forest predicts if the machine will fail by building many decision trees from different parts of the data. It combines their results to make a final call. This cuts down mistakes and gives a solid yes/no on failure, helping avoid surprise breakdowns.

```
Classifier output
Time taken to build model: 1.27 seconds

==== Stratified cross-validation ====
==== Summary ====

Correctly Classified Instances      9854          98.54   %
Incorrectly Classified Instances    146           1.46   %
Kappa statistic                   0.9708
Mean absolute error               0.0408
Root mean squared error          0.112
Relative absolute error           8.1607 %
Root relative squared error     22.4096 %
Total Number of Instances        10000

==== Detailed Accuracy By Class ====

      TP Rate   FP Rate   Precision   Recall   F-Measure   MCC      ROC Area   PRC Area   Class
          0.974     0.004     0.996     0.974     0.985     0.971     0.999     0.999     0
          0.996     0.026     0.975     0.996     0.986     0.971     0.999     0.999     1
Weighted Avg.     0.985     0.015     0.986     0.985     0.985     0.971     0.999     0.999

==== Confusion Matrix ====

      a      b  <-- classified as
4872  128 |    a = 0
  18 4982 |    b = 1
```

The Accuracy: The model correctly classified **98.54%** of instances, meaning it has a high prediction accuracy.

The confusion matrix:

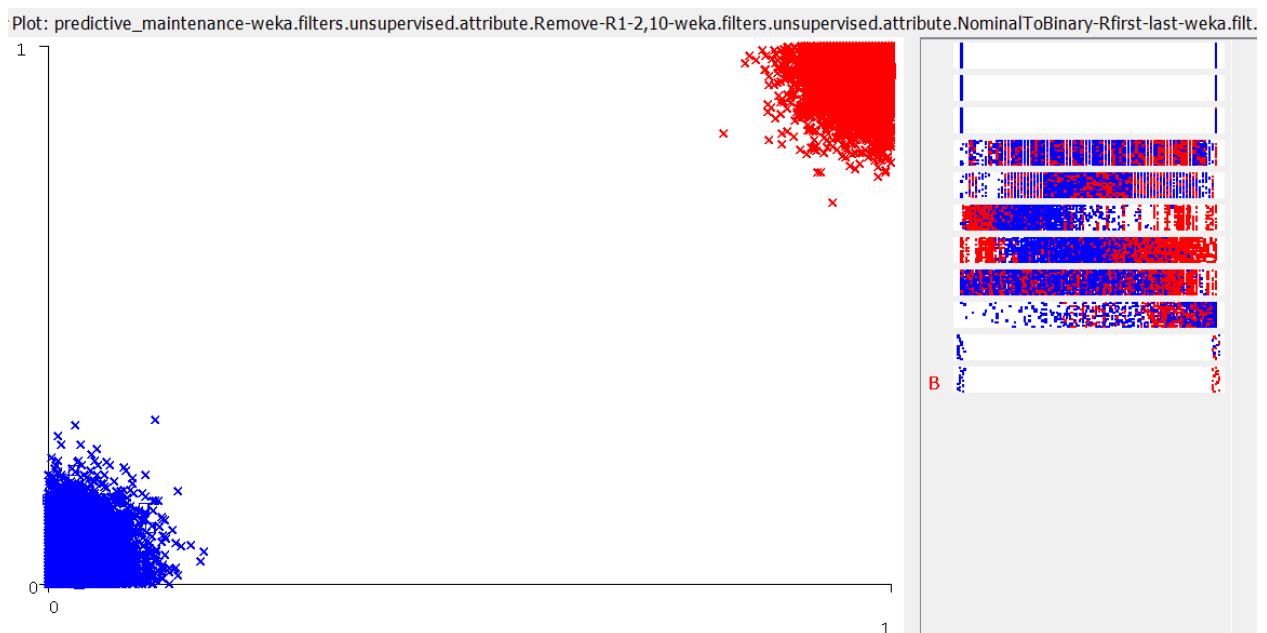
	Predicted 0	Predicted 1
Actual 0	4872	120
Actual 1	18	4902

-The parameter

bagSizePercent	100
batchSize	50
breakTiesRandomly	False
calcOutOfBag	False
computeAttributeImportance	False
debug	False
doNotCheckCapabilities	False
maxDepth	10
numDecimalPlaces	2
numExecutionSlots	1
numFeatures	0
numIterations	100
outputOutOfBagComplexityStatistics	False
printClassifiers	False
seed	1
storeOutOfBagPredictions	False

- I set the Random Forest to build 100 trees with a max depth of 10, using all data for training each tree. We kept the process simple by skipping extra calculations like feature importance and out-of-bag errors, used one CPU core, and fixed the seed for consistent results. This balances accuracy and speed.

-The visualization



15) Naive Bayes

looks at each data feature separately and calculates the chance the machine will fail based on those features. It's fast and works well if features don't depend on each other.

```
Classifier output
Time taken to build model: 0.13 seconds

==== Stratified cross-validation ====
==== Summary ====

Correctly Classified Instances      8938          89.38 %
Incorrectly Classified Instances   1062           10.62 %
Kappa statistic                   0.7876
Mean absolute error               0.1306
Root mean squared error           0.2804
Relative absolute error            26.1158 %
Root relative squared error       56.0796 %
Total Number of Instances         10000

==== Detailed Accuracy By Class ====

      TP Rate  FP Rate  Precision  Recall   F-Measure  MCC     ROC Area  PRC Area  Class
          0.860    0.072    0.923    0.860    0.890    0.789    0.963    0.965      0
          0.928    0.140    0.869    0.928    0.897    0.789    0.963    0.960      1
Weighted Avg.      0.894    0.106    0.896    0.894    0.894    0.789    0.963    0.963

==== Confusion Matrix ====

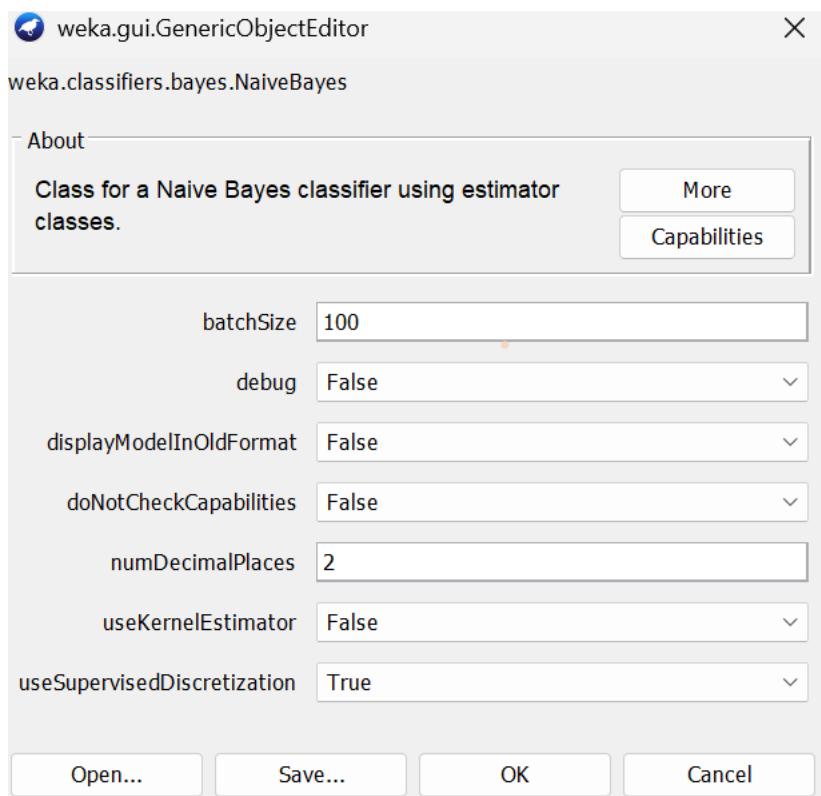
      a     b  <-- classified as
4299  701 |     a = 0
 361 4639 |     b = 1
```

The accuracy: is 89.38%, which is decent but still lower than Random Forest, making it less reliable and too simple for our data.

The confusion matrix:

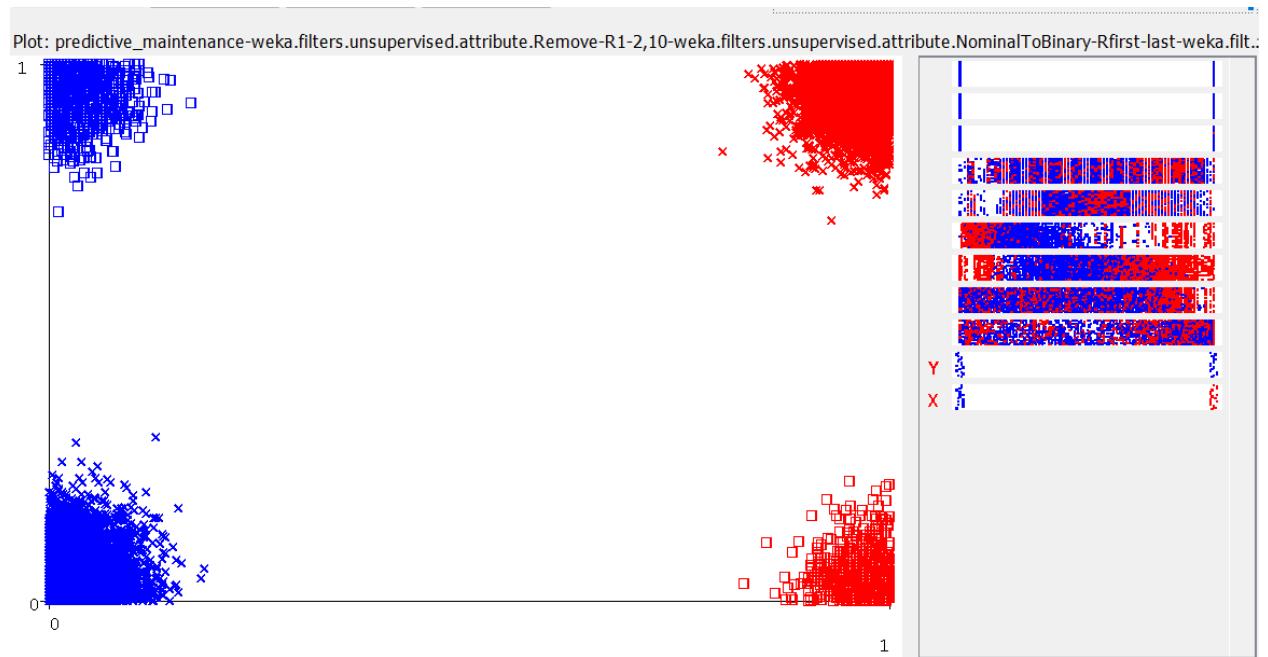
	Predicted 0	Predicted 1
Actual 0	4299	701
Actual 1	361	4636

- The parameter



I kept all parameters at their default settings except for *useSupervisedDiscretization*, which I enabled because it significantly improved accuracy from 82% to 89%. This happens by converting continuous data into class-based discrete categories, helping Naive Bayes perform better.

- The visualization



14) IBk (K-Nearest Neighbors)

is a simple algorithm that decides if a machine will fail by looking at the closest similar machines in the data. It checks the ‘k’ nearest neighbors and picks the majority outcome to predict failure or not. This way, it helps tell apart machines that might fail from those that won’t by comparing them to machines with similar behavior.

```
Classifier output
Time taken to build model: 0 seconds

==== Stratified cross-validation ====
==== Summary ===

Correctly Classified Instances      9692      96.92    %
Incorrectly Classified Instances   308       3.08    %
Kappa statistic                   0.9384
Mean absolute error               0.0312
Root mean squared error          0.1549
Relative absolute error           6.2427 %
Root relative squared error     30.9815 %
Total Number of Instances        10000

==== Detailed Accuracy By Class ===

      TP Rate  FP Rate  Precision  Recall   F-Measure  MCC      ROC Area  PRC Area  Class
0       0.938   0.000    1.000    0.938    0.968    0.940    1.000    1.000     0
1       1.000   0.062    0.942    1.000    0.970    0.940    1.000    1.000     1
Weighted Avg.    0.969   0.031    0.971    0.969    0.969    0.940    1.000    1.000

==== Confusion Matrix ===

      a      b  <-- classified as
4692  308 |  a = 0
  0 5000 |  b = 1
```

The accuracy is 96.92%, which is good but still lower than Random Forest. KNN is considered less reliable and might not fully capture the data’s complexity.

The confusion matrix:

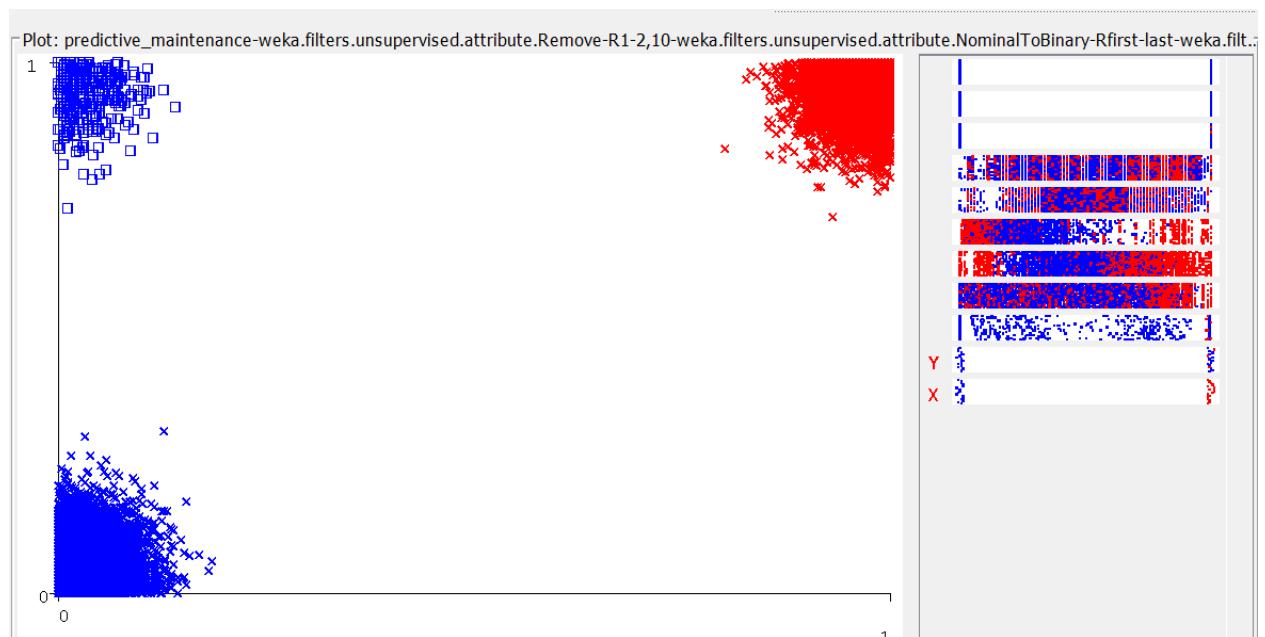
	Predicted 0	Predicted 1
Actual 0	4692	308
Actual 1	0	5000

- The parameter

KNN	6
batchSize	100
crossValidate	False
debug	False
distanceWeighting	Weight by 1/distance
doNotCheckCapabilities	False
meanSquared	False
nearestNeighbourSearchAlgorithm	Choose LinearNNSearch -A "weka...
numDecimalPlaces	2
windowSize	0

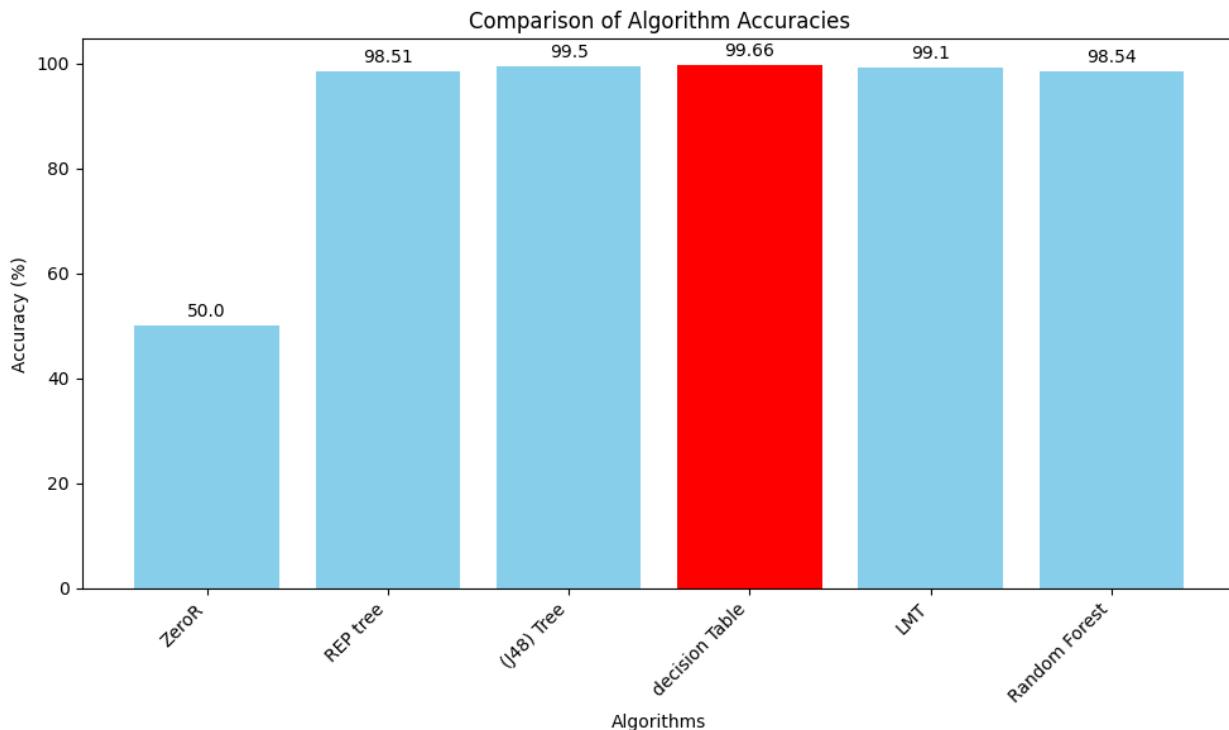
We set KNN to use 6 nearest neighbors and process data in batches of 100 without cross-validation. The model gives more weight to closer neighbors using a simple linear search method. Results are shown with two decimal places, and all neighbors are considered without any window limit. This setup keeps things simple while focusing on the most relevant data points for accurate predictions.

- The visualization



The graph

A comparison of all the algorithms with the best accuracy:



The conclusion:

We tested different algorithms to predict if a machine will fail or not. The best one was the **Decision Table**, with the highest accuracy of **99.66%**. The **J48 Decision Tree** and **AdaBoostM1** also did really well, with accuracy above **99%**.

Other models like **Random Forest** and **REP Tree** worked well too, around **98.5%** accuracy. But simpler models like **Naive Bayes** and **OneR** didn't do as good, which means they might not handle the data complexity well.

The baseline model, **ZeroR**, only got **50%** accuracy, so we need better models to make good predictions.

In short, tree-based models and Decision Table worked best for this data and can help predict machine failures more accurately.

References:

- *About AdaBoostM1* : <https://www.sciencedirect.com/topics/computer-science/adaboost-algorithm>
About Logistic Regression (LR):
- https://youtu.be/J5CM5pdcFgA?si=nhhlnL17WPPz_r8h
About Decision Tree (J48):
- <https://youtu.be/YbDTv8Y4cvI?si=shHWhgWx301EXB0r>
- <https://www.geeksforgeeks.org/building-a-machine-learning-model-using-j48-classifier/>
About OneR :
- <https://christophm.github.io/interpretable-ml-book/rules.html>
About DecisionStump:
- <https://medium.com/geekculture/decision-stump-b8e93c1f54d7>
- *Machine Predictive Maintenance Classification dataset*:
<https://www.kaggle.com/datasets/shivamb/machine-predictive-maintenance-classification?>
- *OpenAI for understanding*:
<https://chatgpt.com/>

